

TRANSACTIONS



Scalaris:

Users and Developers Guide

Version 0.2.0 draft December 28, 2010

Copyright 2007-2010 Konrad-Zuse-Zentrum für Informationstechnik Berlin and onScale solutions.

Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

Contents

I. Users Guide	5
1. Introduction	6
1.1. Brewer's CAP Theorem	6
1.2. Scientific Background	7
2. Download and Installation	8
2.1. Requirements	8
2.2. Download	8
2.2.1. Development Branch	8
2.2.2. Releases	8
2.3. Configuration	8
2.4. Build	9
2.4.1. Linux	9
2.4.2. Windows	9
2.4.3. Java-API	10
2.5. Running Scalaris	10
2.5.1. Running on a local machine	10
2.5.2. Running distributed	11
2.6. Installation	11
2.7. Logging	12
3. Using the system	13
3.1. JSON API	13
3.1.1. Deleting a key	16
3.2. Java command line interface	16
3.3. Java API	17
4. Testing the system	18
4.1. Running the unit tests	18
5. Troubleshooting	19
5.1. Network	19
II. Developers Guide	20
6. General Hints	21
6.1. Coding Guidelines	21
6.2. Testing Your Modifications and Extensions	21
6.3. Help with Digging into the System	21
6.4. General Erlang server loop	21

7. System Infrastructure	23
7.1. Groups of Processes	23
7.2. The Communication Layer <code>comm</code>	23
7.3. The <code>gen_component</code>	23
7.3.1. A basic <code>gen_component</code> including a message handler	24
7.3.2. How to start a <code>gen_component</code> ?	25
7.3.3. When does a <code>gen_component</code> terminate?	25
7.3.4. What happens when unexpected events / messages arrive?	25
7.3.5. What if my message handler generates an exception or crashes the process?	25
7.3.6. Changing message handlers and implementing state dependent message responsiveness as a state-machine	26
7.3.7. Halting and pausing a <code>gen_component</code>	26
7.3.8. Integration with <code>pid_groups</code> : Redirecting events / messages to other <code>gen_components</code>	27
7.3.9. Replying to ping messages	27
7.3.10. The debugging interface of <code>gen_component</code> : Breakpoints and step-wise execution	27
7.3.11. Future use and planned extensions for <code>gen_component</code>	30
7.4. The Process' Database (<code>pdb</code>)	30
7.5. Writing Unittests	30
7.5.1. Plain unittests	30
7.5.2. Randomized Testing using <code>tester.erl</code>	30
8. Basic Structured Overlay	31
8.1. Ring Maintenance	31
8.2. T-Man	31
8.3. Routing Tables	31
8.3.1. The routing table process (<code>rt_loop</code>)	33
8.3.2. Simple routing table (<code>rt_simple</code>)	34
8.3.3. Chord routing table (<code>rt_chord</code>)	37
8.4. Local Datastore	41
8.5. Cyclon	41
8.6. Vivaldi Coordinates	41
8.7. Estimated Global Information (Gossiping)	41
8.8. Load Balancing	41
8.9. Broadcast Trees	41
9. Transactions in Scalaris	42
9.1. The Paxos Module	42
9.2. Transactions using Paxos Commit	42
9.3. Applying the Tx-Modules to replicated DHTs	42
10. How a node joins the system	43
10.1. Supervisor-tree of a Scalaris node	44
10.2. Starting the <code>sup_dht_node</code> supervisor and general processes of a node	44
10.3. Starting the <code>sup_dht_node_core</code> supervisor with a peer and some paxos processes	46
10.4. Initializing a <code>dht_node-process</code>	46
10.5. Actually joining the ring	47
10.5.1. A single node joining an empty ring	47
10.5.2. A single node joining an existing (non-empty) ring	48
11. Directory Structure of the Source Code	56

Part I.

Users Guide

1. Introduction

Scalaris is a scalable, transactional, distributed key-value store based on the peer-to-peer principle. It can be used to build scalable Web 2.0 services. The concept of Scalaris is quite simple: Its architecture consists of three layers.

It provides self-management and scalability by replicating services and data among peers. Without system interruption it scales from a few PCs to thousands of servers. Servers can be added or removed on the fly without any service downtime.



Scalaris takes care of:

- Fail-over
- Data distribution
- Replication
- Strong consistency
- Transactions

The Scalaris project was initiated by Zuse Institute Berlin and onScale solutions and was partly funded by the EU projects Selfman and XtreamOS. Additional information (papers, videos) can be found at <http://www.zib.de/CSR/Projects/scalaris> and <http://www.onscale.de/scalarix.html>.

1.1. Brewer's CAP Theorem

In distributed computing there exists the so called CAP theorem. It basically says that there are three desirable properties for distributed systems but one can only have any two of them.

Strict Consistency. Any read operation has to return the result of the latest write operation on the same data item.

Availability. Items can be read and modified at any time.

Partition Tolerance. The network on which the service is running may split into several partitions which cannot communicate with each other. Later on the networks may re-join again.

For example, a service is hosted on one machine in Seattle and one machine in Berlin. This service is partition tolerant if it can tolerate that all Internet connections over the Atlantic (and Pacific) are interrupted for a few hours and then get repaired.

The goal of Scalaris is to provide strict consistency and partition tolerance. We are willing to sacrifice availability to make sure that the stored data is always consistent. I.e. when you are running Scalaris with a replication degree of 4 and the network splits into two partitions, one partition with three replicas and one partition with one replica, you will be able to continue to use the service only in the larger partition. All requests in the smaller partition will time out until the two networks merge again. Note, most other key-value stores tend to sacrifice consistency.

1.2. Scientific Background

Basics. The general structure of Scalaris is modelled after Chord. The Chord paper [4] describes the ring structure, the routing algorithms, and basic ring maintenance.

The main routines of our Chord node are in `src/dht_node.erl` and the join protocol is implemented in `src/dht_node_join.erl` (see also Chap. [?]). Our implementation of the routing algorithms is described in more detail in Sect. 8.3 and the actual implementation is in `src/rt_chord.erl`.

Transactions. The most interesting part is probably the transaction algorithms. The most current description of the algorithms and background is in [6].

We have currently two generations of transaction algorithms. The older one is in `src/transstore`. The newer one is more modular. It provides an implementation of the paxos algorithm in `src/paxos` and the transaction algorithms itself in `src/transactions` (see also Chap. 9).

Ring Maintenance. We changed the ring maintenance algorithm in Scalaris. It is not the standard Chord one, but a variation of T-Man [5]. It is supposed to fix the ring structure faster. In some situations, the standard Chord algorithm is not able to fix the ring structure while T-Man can still fix it. For node sampling, our implementation relies on Cyclon [7].

The T-Man implementation can be found in `src/rm_tman.erl` and the Cyclon implementation in `src/cyclon`.

Vivaldi Coordinates. For some experiments, we implemented so called Vivaldi coordinates [2]. They can be used to estimate the network latency between arbitrary nodes.

The implementation can be found in `src/vivaldi.erl`.

Gossiping. For some algorithms, we use estimates of global information. These estimates are aggregated with the help of gossiping techniques [8].

The implementation can be found in `src/gossip.erl`.

2. Download and Installation

2.1. Requirements

For building and running Sclaris, some third-party modules are required which are not included in the Sclaris sources:

- Erlang R13B01 or newer
- GNU-like Make

To build the Java API (and the command-line client) the following modules are required additionally:

- Java Development Kit 6
- Apache Ant

Before building the Java API, make sure that `JAVA_HOME` and `ANT_HOME` are set. `JAVA_HOME` has to point to a JDK installation, and `ANT_HOME` has to point to an Ant installation.

2.2. Download

The sources can be obtained from <http://code.google.com/p/scalaris>. RPMs and DEBs are available from <http://download.opensuse.org/repositories/home:/tschuett/>.

2.2.1. Development Branch

You find the latest development version in the svn repository:

```
# Non-members may check out a read-only working copy anonymously over HTTP.  
svn checkout http://scalaris.googlecode.com/svn/trunk/ scalaris-read-only
```

2.2.2. Releases

Releases can be found under the 'Download' tab on the web-page.

2.3. Configuration

Sclaris reads two configuration files from the working directory: `bin/scalaris.cfg` (mandatory) and `bin/scalaris.local.cfg` (optional). The former defines default settings and is included in the release. The latter can be created by the user to alter settings. A sample file is provided as `bin/scalaris.local.cfg.example`. To run Sclaris distributed over several nodes, each node requires a `bin/scalaris.local.cfg`:

File `scalaris.local.cfg`:

```
%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
% Settings for distributed Erlang
% (see scalaris.hrl to switch)

% {boot_host, {boot, 'boot@foo.bar.com'}}.
% {known_hosts, [{service_per_vm, 'boot@foo.bar.com'}]}.

%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
% Settings for TCP mode.
% (see scalaris.hrl to switch)

% Insert the appropriate IP-addresses for your setup
% as comma separated integers:
% IP Address, Port, and label of the boot server
{boot_host, {{127,0,0,1},14195,boot}}.

% IP Address, Port, and label of a node which is already in the system
{known_hosts, [{{{127,0,0,1},14195, service_per_vm}]}].
```

Scalaris currently distinguishes two different kinds of nodes: (a) the boot-server and (b) regular nodes. For the moment, we limit the number of boot-servers to exactly one. The remaining nodes are regular nodes. The boot-server is contacted to join the system. On all servers, the `boot_host` option defines the server where the boot server is running. In the example, it is an IP address plus a TCP port.

2.4. Build

2.4.1. Linux

Scalaris uses `autoconf` for configuring the build environment and `GNU Make` for building the code.

```
%> ./configure
%> make
%> make docs
```

For more details read `README` in the main Scalaris checkout directory.

2.4.2. Windows

We are currently not supporting Scalaris on Windows. However, we have two small bat files for building and running scalaris nodes. It seems to work but we make no guarantees.

- Install Erlang
<http://www.erlang.org/download.html>
- Install OpenSSL (for crypto module)
<http://www.slproweb.com/products/Win32OpenSSL.html>
- Checkout scalaris code from SVN
- adapt the path to your Erlang installation in `build.bat`
- start a `cmd.exe`
- go to the scalaris directory
- run `build.bat` in the cmd window

- check that there were no errors during the compilation; warnings are fine
- go to the bin sub-directory
- adapt the path to your Erlang installation in `boot.bat`, `cs_local.bat`, `cs_local2.bat` and `cs_local3.bat`
- run `boot.bat` or one of the other start scripts in the cmd window

`build.bat` will generate a `Emakefile` if there is none yet. If you have Erlang < R13B04, you will need to adapt the `Emakefile`. There will be empty lines in the first three blocks ending with “`}]`.”: add the following to these lines and try to compile again. It should work now.

```
, {d, type_forward_declarations_are_not_allowed}
, {d, forward_or_recursive_types_are_not_allowed}
```

For the most recent description please see the FAQ at <http://code.google.com/p/scalaris/wiki/FAQ>.

2.4.3. Java-API

The following commands will build the Java API for Scalaris:

```
%> make java
```

This will build `scalaris.jar`, which is the library for accessing the overlay network. Optionally, the documentation can be build:

```
%> cd java-api
%> ant doc
```

2.5. Running Scalaris

As mentioned above, in Scalaris there are two kinds of nodes:

- boot servers
- regular nodes

In every Scalaris, at least one boot server is required. It will maintain a list of nodes taking part in the system and allows other nodes to join the ring. For redundancy, it is also possible to have several boot servers. In the future, we want to eliminate this distinction, so any node is also a boot-server.

2.5.1. Running on a local machine

Open at least two shells. In the first, inside the Scalaris directory, start the boot script (`boot.bat` on Windows):

```
%> ./bin/boot.sh
```

This will start the boot server. On success <http://localhost:8000> should point to the management interface page of the boot server. The main page will show you the number of nodes currently in

the system. After a couple of seconds a first Scalaris should have started in the boot server and the number should increase to one. The main page will also allow you to store and retrieve key-value pairs but should not be used by applications to access Scalaris. See Chapter 3 on page 13 for application APIs.

In a second shell, you can now start a second Scalaris node. This will be a ‘regular server’:

```
%> ./bin/cs_local.sh
```

The second node will read the configuration file and use this information to contact the boot server and join the ring. The number of nodes on the web page should have increased to two by now.

Optionally, a third and fourth node can be started on the same machine. In a third shell:

```
%> ./bin/cs_local2.sh
```

In a fourth shell:

```
%> ./bin/cs_local3.sh
```

This will add 3 nodes to the network. The web pages at <http://localhost:8000> should show the additional nodes.

On linux you can also use the `scalarisctl` script to start boot and ‘regular’ nodes.

2.5.2. Running distributed

Scalaris can be installed on other machines in the same way as described in Section 2.6. In the default configuration, nodes will look for the boot server on `localhost` on port 14195. You should create a `scalaris.local.cfg` pointing to the node running the boot server.

```
% Insert the appropriate IP-addresses for your setup
% as comma separated integers:
% IP Address, Port, and label of the boot server
{boot_host, {{127,0,0,1},14195,boot}}.
```

If you are using the default configuration on the boot server it will listen on port 14195 and you only have to change the IP address in the configuration file. Otherwise the other nodes will not find the boot server. On the remote nodes, you only need to call `./cs_local.sh` and they will automatically contact the configured boot server.

2.6. Installation

For simple tests, you do not need to install Scalaris. You can run it directly from the source directory. Note: `make install` will install scalaris into `/usr/local` and place `scalarisctl` into `/usr/local/bin`. But is more convenient to build an RPM and install it.

```
svn checkout http://scalaris.googlecode.com/svn/trunk/ scalaris-0.0.1
tar -cvjf scalaris-0.0.1.tar.bz2 scalaris-0.0.1 --exclude-vcs
cp scalaris-0.0.1.tar.bz2 /usr/src/packages/SOURCES/
rpmbuild -ba scalaris-0.0.1/contrib/scalaris.spec
```

Your source and binary RPM will be generated in `/usr/src/packages/SRPMS` and `RPMS`. We build RPMs and Debs using checkouts from svn and provide them using the openSUSE BuildService at <http://download.opensuse.org/repositories/home:/tschuett/>. Packages are available for

- Fedora 9, 10, 11, 12, 13,
- Mandriva 2008, 2009, 2009.1, 2010,
- openSUSE 11.0, 11.1, 11.2, 11.3, Factory,
- SLE 10, 11,
- CentOS 5.4,
- RHEL 5,
- Debian 5.0 and
- Ubuntu 9.04, 9.10, 10.04.

Inside those repositories you will also find an erlang package - you don't need this if you already have a recent enough erlang version!

2.7. Logging

Description is based on SVN revision r1083.

Scalaris uses the `log4erl` library (see `contrib/log4erl`) for logging status information and error messages. The log level can be configured in `bin/scalaris.cfg` for both the stdout and file logger. The default value is `warn`; only warnings, errors and severe problems are logged.

```
%% @doc Loglevel: debug < info < warn < error < fatal < none
{log_level, warn}.
{log_level_file, warn}.
```

In some cases, it might be necessary to get more complete logging information, e.g. for debugging. In 10 on page 43, we are explaining the startup process of Scalaris nodes in more detail, here the `info` level provides more detailed information.

```
%% @doc Loglevel: debug < info < warn < error < fatal < none
{log_level, info}.
{log_level_file, info}.
```

3. Using the system

3.1. JSON API

Scalaris supports a JSON API for transactions. To minimize the necessary round trips between a client and Scalaris, it uses request lists, which contain all requests that can be done in parallel. The request list is then send to a Scalaris node with a POST message. The result is an opaque TransLog and a list containing the results of the requests. To add further requests to the transaction, the TransLog and another list of requests may be send to Scalaris. This process may be repeated as often as necessary. To finish the transaction, the request list can contain a 'commit' request as the last element, which triggers the validation phase of the transaction processing.

The JSON-API can be accessed via the Scalaris-Web-Server running on port 8000 by default and the page `jsonrpc.yaws` (For example at: <http://localhost:8000/jsonrpc.yaws>). The following example illustrates the message flow:

Client

Make a transaction, that sets two keys:

Scalaris node

→

```
{
  "method": "req_list",
  "version": "1.1",
  "params":
  [
    [
      { "write": { "keyA": "valueA" } },
      { "write": { "keyB": "valueB" } },
      { "commit": "commit" }
    ]
  ],
  "id": 0
}
```

← Scalaris sends results back

```
{ "result":
  { "results":
    [
      { "op": "commit",
        "value": "ok",
        "key": "ok" },
      { "op": "write",
        "value": "valueB",
        "key": "keyB" },
      { "op": "write",
        "value": "valueA",
        "key": "keyA" }
    ],
    "translog":
      [...]
  },
  "id" : 0
}
```

In a second transaction: Read the two keys →

```
{
  "method": "req_list",
  "version": "1.1",
  "params":
    [
      [
        { "read": "keyA" },
        { "read": "keyB" }
      ]
    ]
  "id": 0
}
```

← Scalaris sends results back

```
{ "result":
  { "results":
    [
      { "op": "read",
        "value": "valueB",
        "key": "keyB" },
      { "op": "read",
        "value": "valueA",
        "key": "keyA" }
    ],
    "translog":
      [...] // this list is the translog
              // for further operations!
              // We name it TLOG here.
  },
  "id" : 0
}
```

Calculate something with the read values →
and make further requests, here a write
and the commit for the whole transaction.
Also include the latest translog we got from
Scalaris (named TLOG here).

```
{
  "method": "req_list",
  "version": "1.1",
  "params":
  [
    TLOG, // translog from prev. result
    [
      { "write": { "keyA": "valueA2" } },
      { "commit": "commit" }
    ]
  ],
  "id" : 0
}
```

← Scalaris sends results back

```
{ "result":
  { "results":
    [ { "op": "commit",
        "value": "ok",
        "key": "ok" },
      { "op": "write",
        "value": "valueA2",
        "key": "keyA" }
    ],
    "translog":
    [...]
  },
  "id" : 0
}
```

A sample usage of the JSON API using Ruby can be found in contrib/jsonrpc.rb.

A single request list must not contain a key more than once!

The allowed requests are:

```
{ "read": "any_key" }
{ "write": { "any_key": "any_value" } }
{ "commit": "commit" }
```

The possible results are:

```
{ "op": "read", "key": "any_key", "value": "any_value" }
{ "op": "read", "key": "any_value", "fail": "reason" } // 'not_found' or 'timeout'

{ "op": "write", "key": "any_key", "value": "any_value" }
{ "op": "read", "key": "any_key", "fail": "reason" }

{ "op": "commit", "value": "ok", "key": "ok" }
{ "op": "commit", "value": "fail", "fail": "reason" }
```

3.1.1. Deleting a key

Outside transactions keys can also be deleted, but it has to be done with care, as explained in the following thread on the mailing list: http://groups.google.com/group/scalaris/browse_thread/thread/ff1d9237e218799.

```
{
  "method": "delete",
  "version": "1.1",
  "params":
    [
      { "key": "any_key" }
    ],
  "id" : 0
}
```

Two sample results

```
{ "result":
  { "ok":2, // how many replicas were deleted successssfully
    "results": [ "ok", "ok", "locks_set", "undef" ]
  }
}
```

```
{ "result":
  { "failure": "reason" }
}
```

3.2. Java command line interface

The jar file contains a small command line interface client. For convenience, we provide a wrapper script called `scalaris` which sets up the Java environment:

```
%> ./java-api/scalaris -help
Script Options:
  --help, --h          print this message and scalaris help
  --noconfig           suppress sourcing of /etc/scalaris/scalaris-java.conf
                      and $HOME/.scalaris/scalaris-java.conf config files
  --execdebug         print scalaris exec line generated by this
                      launch script

usage: scalaris [Options]
  -b,--minibench        run mini benchmark
  -d,--delete <key> <[timeout]> delete an item (default timeout: 2000ms)
                          WARNING: This function can lead to inconsistent data
                          (e.g. deleted items can re-appear). Also when
                          re-creating an item the version before the delete can
                          re-appear.
  -g,--getsubscribers <topic> get subscribers of a topic
  -h,--help             print this message
  -lh,--localhost       gets the local host's name as known to
                          Java (for debugging purposes)
  -p,--publish <topic> <message> publish a new message for the given
                          topic
  -r,--read <key>       read an item
  -s,--subscribe <topic> <url> subscribe to a topic
  -u,--unsubscribe <topic> <url> unsubscribe from a topic
  -v,--verbose          print verbose information, e.g. the
                          properties read
  -w,--write <key> <value> write an item
```


read, write and delete can be used to read, write and delete from/to the overlay, respectively. getsubscribers, publish, and subscribe are the PubSub functions. The others provide debugging and testing functionality.

```
%> ./java-api/scalaris -write foo bar
write(foo, bar)
%> ./java-api/scalaris -read foo
read(foo) == bar
```

Per default, the `scalaris` script tries to connect to a boot server at `localhost`. You can change the node it connects to (and further connection properties) by adapting the values defined in `java-api/scalaris.properties`.

3.3. Java API

The `scalaris.jar` provides the command line client as well as a library for Java programs to access Scalaris. The library provides two classes:

- `Scalaris` provides a high-level API similar to the command line client.
- `Transaction` provides a low-level API to the transaction mechanism.

For details we refer the reader to the Javadoc:

```
%> cd java-api
%> ant doc
%> firefox doc/index.html
```

4. Testing the system

4.1. Running the unit tests

There are some unit tests in the `test` directory. You can call them by running `make test` in the main directory. The results are stored in a local `index.html` file.

The tests are implemented with the `common-test` package from the Erlang system. For running the tests we rely on `run_test`, which is part of the `common-test` package, but (on `erlang < R14`) is not installed by default. `configure` will check whether `run_test` is available. If it is not installed, it will show a warning and a short description of how to install the missing file.

Note: for the unit tests, we are setting up and shutting down several overlay networks. During the shut down phase, the runtime environment will print extensive error messages. These error messages do not indicate that tests failed! Running the complete test suite takes about 3 minutes, depending on your machine. Only if the complete suite finishes, it will present statistics on failed and successful tests.

5. Troubleshooting

5.1. Network

Scalaris uses a couple of TCP ports for communication. It does not use UDP at the moment.

- 8000 HTTP Server on the boot node
- 8001 HTTP Server on the other nodes
- 14195 Port for inter-node communication (boot server)
- 14196 Port for inter-node communication (other nodes)

Please make sure that at least 14195 and 14196 are not blocked by firewalls.

Part II.

Developers Guide

6. General Hints

6.1. Coding Guidelines

- Keep the code short
- Use `gen_component` to implement additional processes
- Don't use `receive` by yourself (Exception: to implement single threaded user API calls (`cs_api`, `yaws_calls`, etc))
- Don't use `erlang:now/0`, `erlang:send_after/3`, `receive after` etc. in performance critical code, consider using `msg_delay` instead.
- Don't use `timer:tc/3` as it catches exceptions. Use `util:tc/3` instead.

6.2. Testing Your Modifications and Extensions

- Run the testsuites using `make test`
- Run the java api test using `make java-test` (Scalaris output will be printed if a test fails; if you want to see it during the tests, start a `bin/boot.sh` and run the tests by `cd java; ant test`)
- Run the Ruby client by starting Scalaris and running `cd contrib; ./jsonrpc.rb`

6.3. Help with Digging into the System

- use `ets:i/0,1` to get details on the local state of some processes
- consider changing `pdb.erl` to use `ets` instead of `erlang:put/get`
- Have a look at `strace -f -p PID` of beam process
- Get message statistics via the Web-interface
- enable/disable tracing for certain modules
- Use `etop` and look at the total memory size and atoms generated
- send processes `sleep` or `kill` messages to test certain behaviour (see `gen_component.erl`)
- use `boot_server:number_of_nodes(). flush().`
- use `admin_checkring(). flush().`

6.4. General Erlang server loop

Servers in Erlang often use the following structure to maintain a state while processing received messages:

```
loop(State) ->  
  receive  
    Message ->  
      State1 = f(State),  
      loop(State1)  
  end.
```

The server runs an endless loop, that waits for a message, processes it and calls itself using tail-recursion in each branch. The loop works on a `State`, which can be modified when a message is handled.

7. System Infrastructure

7.1. Groups of Processes

- What is it? How to distinguish from Erlangs internal named processes?
- Joining a process group
- Why do we do this... (managing several independent nodes inside a single Erlang VM for testing)

7.2. The Communication Layer `comm`

- in general
- format of messages (tuples)
- use messages with cookies (server and client side)
- What is a message tag?

7.3. The `gen_component`

Description is based on SVN revision r993.

The generic component model implemented by `gen_component` allows to add some common functionality to all the components that build up the Scalaris system. It supports:

event-handlers: message handling with a similar syntax as used in [3].

FIFO order of messages: components cannot be inadvertently locked as we do not use selective receive statements in the code.

sleep and halt: for testing components can sleep or be halted.

debugging, breakpoints, stepwise execution: to debug components execution can be steered via breakpoints, step-wise execution and continuation based on arriving events and user defined component state conditions.

basic profiling ,

state dependent message handlers: depending on its state, different message handlers can be used and switched during runtime. Thereby a kind of state-machine based message handling is supported.

prepared for `pid_groups`: allows to send events to named processes inside the same group as the actual component itself (`send_to_group_member`) when just holding a reference to any group member, and

unit-testing of event-handlers: as message handling is separated from the main loop of the component, the handling of individual messages and thereby performed state manipulation can easily be tested in unit-tests by directly calling message handlers.

In *Scalaris* all Erlang processes should be implemented as `gen_component`. The only exception are functions interfacing to the client, where a transition from asynchronous to synchronous request handling is necessary and that are executed in the context of a client's process or a process that behaves as a proxy for a client (`cs_api`).

7.3.1. A basic `gen_component` including a message handler

To implement a `gen_component`, the component has to provide the `gen_component` behaviour:

File `gen_component.erl`:

```
48 -spec behaviour_info(atom()) -> [{atom(), arity()}] | undefined.
49 behaviour_info(callbacks) ->
50 [
51     {init, 1}      % initialize component
52     % note: can use arbitrary on-handler, but by default on/2 is used:
53     {on, 2}        % handle a single message
54     % on(Msg, State) -> NewState | unknown_event | kill
55 ];
```

This is illustrated by the following example:

File `idholder.erl`:

```
72 %% @doc Initialises the idholder with a random key and a counter of 0.
73 -spec init([{{idholder, id}, Id::?RT:key()} | tuple()]) -> state().
74 init(Options) ->
75     Id = case lists:keyfind({idholder, id}, 1, Options) of
76         {{idholder, id}, IdX} -> IdX;
77         _ -> ?RT:get_random_node_id()
78     end,
79     log:log(info, "[ idholder ~w ] init: ~p", [comm:this(), Id]),
80     {Id, 0}.
81
82 -spec on(message(), state()) -> state().
83 on({reinit}, _State) ->
84     {?RT:get_random_node_id(), 0};
85 on({get_id, PID}, {Id, IdVersion} = State) ->
86     comm:send_local(PID, {idholder_get_id_response, Id, IdVersion}),
87     State;
88 on({set_id, NewId, NewIdVersion}, _State) ->
89     {NewId, NewIdVersion};
90 on({web_debug_info, Requestor}, {Id, IdVersion} = State) ->
91     KeyValueCollection =
92         [{"id", lists:flatten(io_lib:format("~p", [Id]))},
93          {"id_version", lists:flatten(io_lib:format("~p", [IdVersion]))}],
94     comm:send_local(Requestor, {web_debug_info_reply, KeyValueCollection}),
95     State.
```

`your_gen_component:init/1` is called during start-up of a `gen_component` and should return the initial state to be used for this `gen_component`. Later, the current state of the component can be retrieved using `gen_component:get_state/1`.

To react on messages / events, a message handler is used. The default message handler is called `your_gen_component:on/2`. This can be changed by calling `gen_component:change_handler/2` (see Section 7.3.6). When an event / message for the component arrives, this handler is called with the event itself and the current state of the component. In the handler, the state of the component may be adjusted depending upon the event. The handler itself may trigger new events / messages for itself or other components and has finally to return the updated state of the component or the atoms `unknown_event` or `kill`. It must neither call `receive` nor `timer:sleep/1` nor `erlang:exit/1`.

7.3.2. How to start a gen_component?

A `gen_component` can be started using one of:

```
gen_component:start(Module, Args, GenCOptions = [])
```

```
gen_component:start_link(Module, Args, GenCOptions = [])
```

Module: the name of the module your component is implemented in

Args: List of parameters passed to `Module:init/1` for initialization

GenCOptions: optional parameter. List of options for `gen_component`

`{pid_groups_join_as, ProcessGroup, ProcessName}`: registers the new process with the given process group (also called instanceid) and name using `pid_groups`.

`{erlang_register, ProcessName}`: registers the process as a named Erlang process.

`wait_for_init`: wait for `Module:init/1` to return before returning to the caller.

These functions are compatible to the Erlang/OTP supervisors. They spawn a new process for the component which itself calls `Module:init/1` with the given `Args` to initialize the component. `Module:init/1` should return the initial state for your component. For each message sent to this component, the default message handler `Module:on(Message, State)` will be called, which should react on the message and return the updated state of your component.

`gen_component:start()` and `gen_component:start_link()` return the pid of the spawned process as `{ok, Pid}`.

7.3.3. When does a gen_component terminate?

A `gen_component` can be stopped using:

`gen_component:kill(Pid)` or by returning `kill` from the current message handler.

7.3.4. What happens when unexpected events / messages arrive?

Your message handler (default is `your_gen_component:on/2`) should return `unknown_event` in the final clause (`your_gen_component:on(_, _)`). `gen_component` then will nicely report on the unhandled message, the component's name, its state and currently active message handler, as shown in the following example:

```
# bin/boot.sh
[...]
(boot@localhost)10> pid_groups ! {no_message}.
{no_message}
[error] unknown message: {no_message} in Module: pid_groups and
handler on in State null
(boot@localhost)11>
```

The `pid_groups` (see Section 7.1) is a `gen_component` which registers itself as named Erlang process with the `gen_component` option `erlang_register` and therefore can be addressed by its name in the Erlang shell. We send it a `{no_message}` and `gen_component` reports on the unhandled message. The `pid_groups` module itself continues to run and waits for further messages.

7.3.5. What if my message handler generates an exception or crashes the process?

`gen_component` catches exceptions generated by message handlers and reports them with a stack trace, the message, that generated the exception, and the current state of the component.

If a message handler terminates the process via `erlang:exit/1`, this is out of the responsibility scope of `gen_component`. As usual in Erlang, all linked processes will be informed. If for example `gen_component:start_link/2` or `/3` was used for starting the `gen_component`, the spawning process will be informed, which may be an Erlang supervisor process taking further actions.

7.3.6. Changing message handlers and implementing state dependent message responsiveness as a state-machine

Sometimes it is beneficial to handle messages depending on the state of a component. One possibility to express this is implementing different clauses depending on the state variable, another is introducing case clauses inside message handlers to distinguish between current states. Both approaches may become tedious, error prone, and may result in confusing source code.

Sometimes the use of several different message handlers for different states of the component leads to clearer arranged code, especially if the set of handled messages changes from state to state. For example, if we have a component with an initialization phase and a production phase afterwards, we can handle in the first message handler messages relevant during the initialization phase and simply queue all other requests for later processing using a common default clause.

When initialization is done, we handle the queued user requests and switch to the message handler for the production phase. The message handler for the initialization phase does not need to know about messages occurring during production phase and the message handler for the production phase does not need to care about messages used during initialization. Both handlers can be made independent and may be extended later on without any adjustments to the other.

One can also use this scheme to implement complex state-machines by changing the message handler from state to state.

To switch the message handler `gen_component:change_handler(State, new_handler)` is called as the last operation after a message in the active message handler was handled, so that the return value of `gen_component:change_handler/2` is propagated to `gen_component`. The new handler is given as an atom, which is the name of the 2-ary function in your component module to be called.

Starting with non-default message handler.

It is also possible to change the message handler right from the start in your `your_gen_component:init/1` to avoid the default message handler `your_gen_component:on/2`. Just create your initial state as usual and call `gen_component:change_handler(State, my_handler)` as the final call in your `your_gen_component:init/1`. We prepared `gen_component:change_handler/2` to return `State` itself, so this will work properly.

7.3.7. Halting and pausing a `gen_component`

Using `gen_component:kill(Pid)` and `gen_component:sleep(Pid, Time)` components can be terminated or paused.

7.3.8. Integration with `pid_groups`: Redirecting events / messages to other `gen_components`

Each `gen_component` by itself is prepared to support `comm:send_to_group_member/3` which forwards messages inside a group of processes registered via `pid_groups` (see Section 7.1) by their name. So, if you hold a `Pid` of one member of a process group, you can send messages to other members of this group, if you know their registered Erlang name. You do not necessarily have to know their individual `Pid`.

In consequence, no `gen_component` can individually handle messages of the form `{send_to_group_member, _, _}` as such messages are consumed by `gen_component` itself.

7.3.9. Replying to `ping` messages

Each `gen_component` replies automatically to `{ping, Pid}` requests with a `{pong}` send to the given `Pid`. Such messages are generated, for example, by `vivaldi_latency` which is used by our `vivaldi` module.

In consequence, no `gen_component` can individually handle messages of the form: `{ping, _}` as such messages are consumed by `gen_component` itself.

7.3.10. The debugging interface of `gen_component`: Breakpoints and step-wise execution

We equipped `gen_component` with a debugging interface, which especially is beneficial, when testing the interplay between several `gen_components`. It supports breakpoints (bp) which can pause the `gen_component` depending on the arriving messages or depending on user defined conditions. If a breakpoint is reached, the execution can be continued step-wise (message by message) or until the next breakpoint is reached.

We use it in our unit tests to steer protocol interleavings and to perform tests using random protocol interleavings between several processes (see `paxos_SUITE`). It allows also to reproduce given protocol interleavings for better testing.

Managing breakpoints.

Breakpoints are managed by the following functions:

`gen_component:bp_set(Pid, MsgTag, BPName)`: For the component running under `Pid` a breakpoint `BPName` is set. It is reached, when a message with a message tag `MsgTag` is next to be handled by the component (See `comm:get_msg_tag/1` and Section 7.2 for more information on message tags). The `BPName` is used as a reference for this breakpoint, for example to delete it later.

`gen_component:bp_set_cond(Pid, Cond, BPName)`: The same as `gen_component:bp_set/3` but a user defined condition implemented in `{Module, Function, Params = 2}` = `Cond` is checked by calling `Module:Function(Message, State)` to decide whether a breakpoint is reached or not. `Message` is the next message to be handled by the component and `State` is the current state of the component. `Module:Function/2` should return a boolean.

`gen_component:bp_del(Pid, BPName)`: The breakpoint `BPName` is deleted. If the component is in this breakpoint, it will not be released by this call. This has to be done separately by

`gen_component:bp_cont/1`. But the deleted breakpoint will no longer be considered for newly entering a breakpoint.

`gen_component:bp_barrier(Pid)`: Delay all further handling of breakpoint requests until a breakpoint is actually entered.

Note, that the following call sequence may not catch the breakpoint at all, as during the sleep the component not necessarily consumes a ping message and the set breakpoint 'sample_bp' may already be deleted before a ping message arrives.

```
gen_component:bp_set(Pid, ping, sample_bp),
timer:sleep(10),
gen_component:bp_del(Pid, sample_bp),
gen_component:bp_cont(Pid).
```

To overcome this, `gen_component:bp_barrier/1` can be used:

```
gen_component:bp_set(Pid, ping, sample_bp),
gen_component:bp_barrier(Pid),
%% After the bp_barrier request, following breakpoint requests
%% will not be handled before a breakpoint is actually entered.
%% The gen_component itself is still active and handles messages as usual
%% until it enters a breakpoint.
gen_component:bp_del(Pid, sample_bp),
% Delete the breakpoint after it was entered once (ensured by bp_barrier).
% Release the gen_component from the breakpoint and continue.
gen_component:bp_cont(Pid).
```

None of the calls in the sample listing above is blocking. It just schedules all the operations, including the `bp_barrier`, for the `gen_component` and immediately finishes. The actual events of entering and continuing the breakpoint in the `gen_component` happens independently later on, when the next ping message arrives.

Managing execution.

The execution of a `gen_component` can be managed by the following functions:

`gen_component:bp_step(Pid)`: This is the only blocking breakpoint function. It waits until the `gen_component` is in a breakpoint and has handled a single message. It returns the module, the active message handler, and the handled message as a tuple `{Module, On, Message}`. This function does not actually finish the breakpoint, but just lets a single message pass through. For further messages, no breakpoint condition has to be valid, the original breakpoint is still active. To leave a breakpoint, use `gen_component:bp_cont/1`.

`gen_component:bp_cont(Pid)`: Leaves a breakpoint. `gen_component` runs as usual until the next breakpoint is reached.

If no further breakpoints should be entered after continuation, you should delete the registered breakpoint using `gen_component:bp_del/2` before continuing the execution with `gen_component:bp_cont/1`. To ensure, that the breakpoint is entered at least once, `gen_component:bp_barrier/1` should be used before deleting the breakpoint (see the example above). Otherwise it could happen, that the delete request arrives at your `gen_component` before it was actually triggered. The following continuation request would then unintentional apply to an unrelated breakpoint that may be entered later on.

`gen_component:runnable(Pid)`: Returns whether a `gen_component` has messages to handle and is runnable. If you know, that a `gen_component` is in a breakpoint, you can use this to check,

whether a `gen_component:bp_step/1` or `gen_component:bp_cont/1` is applicable to the component.

Tracing handled messages – getting a message interleaving protocol.

We use the debugging interface of `gen_component` to test protocols with random interleaving. First we start all the components involved, set breakpoints on the initialization messages for a new Paxos consensus and then start a single Paxos instance on all of them. The outcome of the Paxos consensus is a `learner_decide` message. So, in `paxos_SUITE:step_until_decide/3` we look for runnable processes and select randomly one of them to perform a single step until the protocol finishes with a decision.

File `paxos_SUITE.erl`:

```
222 -spec(prop_rnd_interleave/3 :: (1..4, 4..16, {pos_integer(), pos_integer(), pos_integer()}
223     -> boolean()).
224 prop_rnd_interleave(NumProposers, NumAcceptors, Seed) ->
225   ct:pal("Called with: paxos_SUITE:prop_rnd_interleave(~p, ~p, ~p).~n",
226     [NumProposers, NumAcceptors, Seed]),
227   Majority = NumAcceptors div 2 + 1,
228   {Proposers, Acceptors, Learners} =
229     make(NumProposers, NumAcceptors, 1, rnd_interleave),
230   %% set bp on all processes
231   [ gen_component:bp_set(element(3, X), proposer_initialize, bp)
232     || X <- Proposers ],
233   [ gen_component:bp_set(element(3, X), acceptor_initialize, bp)
234     || X <- Acceptors ],
235   [ gen_component:bp_set(element(3, X), learner_initialize, bp)
236     || X <- Learners ],
237   %% start paxos instances
238   [ proposer:start_paxosid(X, paxidrndinterl, Acceptors,
239     proposal, Majority, NumProposers, Y)
240     || {X,Y} <- lists:zip(Proposers, lists:seq(1, NumProposers)) ],
241   [ acceptor:start_paxosid(X, paxidrndinterl, Learners)
242     || X <- Acceptors ],
243   [ learner:start_paxosid(X, paxidrndinterl, Majority,
244     comm:this(), cpaxidrndinterl)
245     || X <- Learners ],
246   %% randomly step through protocol
247   OldSeed = random:seed(Seed),
248   Steps = step_until_decide(Proposers ++ Acceptors ++ Learners, cpaxidrndinterl, 0),
249   ct:pal("Needed ~p steps~n", [Steps]),
250   case OldSeed of
251     undefined -> ok;
252     _ -> random:seed(OldSeed)
253   end,
254   true.
255
256 step_until_decide(Processes, PaxId, SumSteps) ->
257   %% io:format("Step ~p~n", [SumSteps]),
258   Runnable = [ X || X <- Processes, gen_component:runnable(element(3,X)) ],
259   case Runnable of
260     [] ->
261       ct:pal("No runnable processes of ~p~n", [length(Processes)]),
262       timer:sleep(5), step_until_decide(Processes, PaxId, SumSteps);
263     _ -> ok
264   end,
265   Num = random:uniform(length(Runnable)),
266   gen_component:bp_step(element(3,lists:nth(Num, Runnable))),
267   receive
268     {learner_decide, cpaxidrndinterl, _, _Res} = _Any ->
269       %% io:format("Received ~p~n", [_Any]),
270       SumSteps
271   after 0 -> step_until_decide(Processes, PaxId, SumSteps + 1)
272   end.
```

To get a message interleaving protocol, we either can output the results of each `gen_component:-bp_step/1` call together with the `Pid` we selected for stepping, or alter the definition of the macro `TRACE_BP_STEPS` in `gen_component`, when we execute all `gen_components` locally in the same Erlang virtual machine.

File `gen_component.erl`:

```
31 %-define(TRACE_BP_STEPS(X,Y), io:format(X,Y)). %% output on console  
32 %-define(TRACE_BP_STEPS(X,Y), ct:pal(X,Y)). %% output even if called by unittest  
33 -define(TRACE_BP_STEPS(X,Y), ok).
```

7.3.11. Future use and planned extensions for `gen_component`

`gen_component` could be further extended. For example it could support hot-code upgrade or could be used to implement algorithms that have to be run across several components of `Scalaris` like snapshot algorithms or similar extensions.

7.4. The Process' Database (pdb)

- How to use it and how to switch from `erlang:put/set` to `ets` and implied limitations.

7.5. Writing Unittests

7.5.1. Plain unittests

7.5.2. Randomized Testing using `tester.erl`

8. Basic Structured Overlay

8.1. Ring Maintenance

8.2. T-Man

8.3. Routing Tables

Description is based on SVN revision r1236.

Each node of the ring can perform searches in the overlay.

A search is done by a lookup in the overlay, but there are several other demands for communication between peers. Scalaris provides a general interface to route a message to the (other) peer, which is currently responsible for a given key.

File `lookup.erl`:

```
30 -spec unreliable_lookup(Key::?RT:key(), Msg::comm:message()) -> ok.
31 unreliable_lookup(Key, Msg) ->
32     comm:send_local(pid_groups:find_a(dht_node),
33                     {lookup_aux, Key, 0, Msg}).
34
35 -spec unreliable_get_key(Key::?RT:key()) -> ok.
36 unreliable_get_key(Key) ->
37     unreliable_lookup(Key, {get_key, comm:this(), Key}).
38
39 -spec unreliable_get_key(CollectorPid::comm:mypid(),
40                         ReqId::{rdht_req_id, pos_integer()},
41                         Key::?RT:key()) -> ok.
42 unreliable_get_key(CollectorPid, ReqId, Key) ->
43     unreliable_lookup(Key, {get_key, CollectorPid, ReqId, Key}).
```

The message `Msg` could be a `get_key` which retrieves content from the responsible node or a `get_node` message, which returns a pointer to the node.

All currently supported messages are listed in the file `dht_node.erl`.

The message routing is implemented in `dht_node_lookup.erl`

File `dht_node_lookup.erl`:

```
27 %% @doc Find the node responsible for Key and send him the message Msg.
28 -spec lookup_aux(State::dht_node_state:state(), Key::intervals:key(),
29                 Hops::non_neg_integer(), Msg::comm:message()) -> ok.
30 lookup_aux(State, Key, Hops, Msg) ->
31     case intervals:in(Key, dht_node_state:get(State, succ_range)) of
32     true -> % found node -> terminate
33         P = dht_node_state:get(State, succ_pid),
34         comm:send(P, {lookup_fin, Key, Hops + 1, Msg});
35     _ ->
36         P = ?RT:next_hop(State, Key),
37         comm:send(P, {lookup_aux, Key, Hops + 1, Msg})
38     end.
```

Each node is responsible for a certain key interval. The function `intervals:in/2` is used to decide, whether the key is between the current node and its successor. If that is the case, the final step is delivers a `lookup_fin` message to the local node. Otherwise, the message is forwarded to the next nearest known peer (listed in the routing table) determined by `?RT:next_hop/2`.

`rt_beh.erl` is a generic interface for routing tables. It can be compared to interfaces in Java. In Erlang interfaces can be defined using a so called ‘behaviour’. The files `rt_simple` and `rt_chord` implement the behaviour ‘`rt_beh`’.

The macro `?RT` is used to select the current implementation of routing tables. It is defined in `include/scalaris.hrl`.

File `scalaris.hrl`:

```

25 %%The RT macro determines which kind of routingtable is used. Uncomment the
26 %%one that is desired.
27
28 %%Standard Chord routingtable
29 -define(RT, rt_chord).
30 -define(MINUS_INFINITY, 0).
31 -define(PLUS_INFINITY, 16#FFFFFFFFFFFFFFFFFFFFFFFFFFFFFFFF).
32
33 %%Simple routingtable
34 %-define(RT, rt_simple).

```

The functions, that have to be implemented for a routing mechanism are defined in the following file:

File `rt_beh.erl`:

```

32 -spec behaviour_info(atom()) -> [{atom(), arity()}] | undefined.
33 behaviour_info(callbacks) ->
34 [
35     % create a default routing table
36     {empty, 1}, {empty_ext, 1},
37     % mapping: key space -> identifier space
38     {hash_key, 1}, {get_random_node_id, 0},
39     % routing
40     {next_hop, 2},
41     % trigger for new stabilization round
42     {init_stabilize, 2},
43     % adapt RT to changed neighborhood
44     {update, 3},
45     % dead nodes filtering
46     {filter_dead_node, 2},
47     % statistics
48     {to_pid_list, 1}, {get_size, 1},
49     % gets all (replicated) keys for a given (hashed) key
50     % (for symmetric replication)
51     {get_replica_keys, 1},
52     % address space size, range and split key
53     % (may all throw 'throw:not_supported' if unsupported by the RT)
54     {n, 0}, {get_range, 2}, {get_split_key, 2},
55     % for debugging and web interface
56     {dump, 1},
57     % for bulkowner
58     {to_list, 1},
59     % convert from internal representation to version for dht_node
60     {export_rt_to_dht_node, 2},
61     % handle messages specific to a certain routing-table implementation
62     {handle_custom_message, 2},
63     % common methods
64     {check, 4}, {check, 5},
65     {check_config, 0}
66 ];

```


empty/1 gets a successor and generates an empty routing table for use inside the routing table implementation. The data structure of the routing table is undefined. It can be a list, a tree, a matrix ...

empty_ext/1 similarly creates an empty external routing table for use by the dht_node. This process might not need all the information a routing table implementation requires and can thus work with less data.

hash_key/1 gets a key and maps it into the overlay's identifier space.

get_random_node_id/0 returns a random node id from the overlay's identifier space. This is used for example when a new node joins the system.

next_hop/2 gets a dht_node's state (including the external routing table representation) and a key and returns the node, that should be contacted next when searching for the key, i.e. the known node nearest to the id.

init_stabilize/2 is called periodically to rebuild the routing table. The parameters are the identifier of the node, its successor and the old (internal) routing table state. This method may send messages to the routing_table process which need to be handled by the handle_custom_message/2 handler since they are implementation-specific.

update/7 is called when the node's ID, predecessor and/or successor changes. It updates the (internal) routing table with the (new) information.

filter_dead_node/2 is called by the failure detector and tells the routing table about dead nodes. This function gets the (internal) routing table and a node to remove from it. A new routing table state is returned.

to_pid_list/1 get the PIDs of all (internal) routing table entries.

get_size/1 get the (internal or external) routing table's size.

get_replica_keys/1 Returns for a given (hashed) Key the (hashed) keys of its replicas. This used for implementing symmetric replication.

n/0 gets the number of available keys. An implementation may throw `throw:not_supported` if the operation is unsupported by the routing table.

dump/1 dump the (internal) routing table state for debugging, e.g. by using the web interface. Returns a list of `{Index, Node_as_String}` tuples which may just as well be empty.

to_list/1 convert the (external) representation of the routing table inside a given `dht_node_state` to a sorted list of known nodes from the routing table, i.e. first=succ, second=next known node on the ring, ... This is used by bulk-operations to create a broadcast tree.

export_rt_to_dht_node/2 convert the internal routing table state to an external state. Gets the internal state and the node's neighborhood for doing so.

handle_custom_message/2 handle messages specific to the routing table implementation. `rt_loop` will forward unknown messages to this function.

check/5, check/6 check for routing table changes and send an updated (external) routing table to the dht_node process.

check_config/0 check that all required configuration parameters exist and satisfy certain restrictions.

8.3.1. The routing table process (rt_loop)

The `rt_loop` module implements the process for all routing tables. It processes messages and calls the appropriate methods in the specific routing table implementations.

File `rt_loop.erl`:

```

40 -opaque(state_active() :: {NeighbTable :: tid(),
41                               RTState    :: ?RT:rt(),
42                               TriggerState :: trigger:state()}).
```

```

43 -type(state_inactive() :: {inactive,
44                               MessageQueue::msg_queue:msg_queue(),
45                               TriggerState::trigger:state()}).
46 %% -type(state() :: state_active() | state_inactive()).

```

If initialized, the node's id, its predecessor, successor and the routing table state of the selected implementation (the macro RT refers to).

File `rt_loop.erl`:

```

154 on_active({trigger_rt}, {NeighbTable, OldRT, TriggerState}) ->
155     % start periodic stabilization
156     % log:log(debug, "[ RT ] stabilize"),
157     Neighbors = rm_loop:get_neighbors(NeighbTable),
158     NewRT = ?RT:init_stabilize(Neighbors, OldRT),
159     ?RT:check(OldRT, NewRT, Neighbors, true),
160     % trigger next stabilization
161     NewTriggerState = trigger:next(TriggerState),
162     new_state(NeighbTable, NewRT, NewTriggerState);

```

Periodically (see `routingtable_trigger` and `pointer_base_stabilization_interval` config parameters) a trigger message is sent to the `rt_loop` process that starts the periodic stabilization implemented by each routing table.

File `rt_loop.erl`:

```

138 % update routing table with changed ID, pred and/or succ
139 on_active({update_rt, OldNeighbors}, {NeighbTable, OldRT, TriggerState}) ->
140     NewNeighbors = rm_loop:get_neighbors(NeighbTable),
141     case ?RT:update(OldRT, OldNeighbors, NewNeighbors) of
142         {trigger_rebuild, NewRT} ->
143             % trigger immediate rebuild
144             NewTriggerState = trigger:now(TriggerState),
145             ?RT:check(OldRT, NewRT, OldNeighbors, NewNeighbors, true),
146             new_state(NeighbTable, NewRT, NewTriggerState);
147         {ok, NewRT} ->
148             ?RT:check(OldRT, NewRT, OldNeighbors, NewNeighbors, true),
149             new_state(NeighbTable, NewRT, TriggerState)
150     end;

```

Every time a node's neighborhood changes, the `dht_node` sends an `update_rt` message to the routing table which will call `?RT:update/7` that decides whether the routing table should be re-build. If so, it will stop any waiting trigger and schedule an immediate (periodic) stabilization.

8.3.2. Simple routing table (`rt_simple`)

One implementation of a routing table is the `rt_simple`, which routes via the successor. Note that this is inefficient as it needs a linear number of hops to reach its goal. A more robust implementation, would use a successor list. This implementation is also not very efficient in the presence of churn.

Data types

First, the data structure of the routing table is defined:

File `rt_simple.erl`:

```

26 -type key_t() :: 0..16#FFFFFFFFFFFFFFFFFFFFFFFFFFFFFFFF. % 128 bit numbers
27 -type rt_t() :: Succ::node:node_type().

```

```

28 -type external_rt_t() :: Succ::node:node_type().
29 -type custom_message() :: none().

```

The routing table only consists of a node (the successor). Keys in the overlay are identified by integers ≥ 0 .

A simple `rm_beh` behaviour

File `rt_simple.erl`:

```

41 %% @doc Creates an "empty" routing table containing the successor.
42 empty(Neighbors) -> nodelist:succ(Neighbors).

```

File `rt_simple.erl`:

```

206 empty_ext(Neighbors) -> empty(Neighbors).

```

The empty routing table (internal or external) consists of the successor.

File `rt_simple.erl`:

Keys are hashed using MD5 and have a length of 128 bits.

File `rt_simple.erl`:

```

61 %% @doc Generates a random node id, i.e. a random 128-bit number.
62 get_random_node_id() ->
63     case config:read(key_creator) of
64     random -> hash_key_(randoms:getRandomId());
65     random_with_bit_mask ->
66         {Mask1, Mask2} = config:read(key_creator_bitmask),
67         (hash_key_(randoms:getRandomId()) band Mask2) bor Mask1
68     end.

```

Random node id generation uses the helpers provided by the `randoms` module.

File `rt_simple.erl`:

```

210 %% @doc Returns the next hop to contact for a lookup.
211 next_hop(State, _Key) -> node:pidX(dht_node_state:get(State, rt)).

```

Next hop is always the successor.

File `rt_simple.erl`:

```

76 %% @doc Triggered by a new stabilization round, renews the routing table.
77 init_stabilize(Neighbors, _RT) -> empty(Neighbors).

```

`init_stabilize/2` resets its routing table to the current successor.

File `rt_simple.erl`:

```

81 %% @doc Updates the routing table due to a changed node ID, pred and/or succ.
82 -spec update(OldRT::rt(), OldNeighbors::nodelist:neighborhood(),
83             NewNeighbors::nodelist:neighborhood()) -> {ok, rt()}.
84 update(_OldRT, _OldNeighbors, NewNeighbors) ->
85     {ok, nodelist:succ(NewNeighbors)}.

```

update/7 updates the routing table with the new successor.

File `rt_simple.erl`:

```

89  %% @doc Removes dead nodes from the routing table (rely on periodic
90  %%      stabilization here).
91  filter_dead_node(RT, _DeadPid) -> RT.

```

`filter_dead_node/2` does nothing, as only the successor is listed in the routing table and that is reset periodically in `init_stabilize/2`.

File `rt_simple.erl`:

```
95 %% @doc Returns the pids of the routing table entries.
96 to_pid_list(Succ) -> [node:pidX(Succ)].
```

to_pid_list/1 returns the pid of the successor.

File rt_simple.erl:

```
100 %% @doc Returns the size of the routing table.
101 get_size(_RT) -> 1.
```

The size of the routing table is always 1.

File rt_simple.erl:

```

138 %% @doc Returns the replicas of the given key.
139 get_replica_keys(Key) ->
140     [Key,
141      Key bxor 16#40000000000000000000000000000000,
142      Key bxor 16#80000000000000000000000000000000,
143      Key bxor 16#C0000000000000000000000000000000
144     ].

```

This `get_replica_keys/1` implements symmetric replication.

File `rt_simple.erl`:[illegible]

There are 2^{128} available keys.

File `rt_simple.erl`:

```
148 %% @doc Dumps the RT state for output in the web interface.
149 dump(Succ) -> [{"0", lists:flatten(io_lib:format("~p", [Succ]))}].
```

dump/1 lists the successor.

File `rt_simple.erl`:

```

221 % @doc Converts the (external) representation of the routing table to a list
222 %     in the order of the fingers, i.e. first=succ, second=shortest finger,
223 %     third=next longer finger,...
224 to_list(State) -> [dht_node_state:get(State, rt)].

```

to_list/1 lists the successor from the external routing table state.

File `rt_simple.erl`:

```
215 %% @doc Converts the internal RT to the external RT used by the dht_node. Both
216 %%     are the same here.
217 export_rt_to_dht_node(RT, _Neighbors) -> RT.
```

`export_rt_to_dht_node/2` states that the external routing table is the same as the internal table.

File `rt_simple.erl`:

```
167 %% @doc There are no custom messages here.
168 -spec handle_custom_message
169     (custom_message() | any(), rt_loop:state_active()) -> unknown_event.
170 handle_custom_message(_Message, _State) -> unknown_event.
```

Custom messages could be send from a routing table process on one node to the routing table process on another node and are independent from any other implementation.

File `rt_simple.hrl`:

```
174 %% @doc Notifies the dht_node and failure detector if the routing table changed.
175 %%     Provided for convenience (see check/5).
176 check(OldRT, NewRT, Neighbors, ReportToFD) ->
177     check(OldRT, NewRT, Neighbors, Neighbors, ReportToFD).
178
179 %% @doc Notifies the dht_node if the (external) routing table changed.
180 %%     Also updates the failure detector if ReportToFD is set.
181 %%     Note: the external routing table only changes the internal RT has
182 %%     changed.
183 check(OldRT, NewRT, _OldNeighbors, NewNeighbors, ReportToFD) ->
184     case OldRT == NewRT of
185         true -> ok;
186         _ ->
187             Pid = pid_groups:get_my(dht_node),
188             RT_ext = export_rt_to_dht_node(NewRT, NewNeighbors),
189             comm:send_local(Pid, {rt_update, RT_ext}),
190             % update failure detector:
191             case ReportToFD of
192                 true ->
193                     NewPids = to_pid_list(NewRT),
194                     OldPids = to_pid_list(OldRT),
195                     fd:update_subscriptions(OldPids, NewPids);
196                 _ -> ok
197             end
198     end.
```

Checks whether the routing table changed and in this case sends the `dht_node` an updated (external) routing table state. Optionally the failure detector is updated. This may not be necessary, e.g. if `check` is called after a crashed node has been reported by the failure detector (the failure detector already unsubscribes the node in this case).

8.3.3. Chord routing table (`rt_chord`)

The file `rt_chord.erl` implements Chord's routing.

Data types

File `rt_chord.erl`:

```
26 -type key_t() :: 0..16#FFFFFFFFFFFFFFFFFFFFFFFFFFFFFFFF. % 128 bit numbers
27 -type rt_t() :: gb_tree().
```

```

28 -type external_rt_t() :: gb_tree().
29 -type index() :: {pos_integer(), non_neg_integer()}.
30 -opaque custom_message() ::
31     {rt_get_node, Source_PID::comm:mypid(), Index::index()} |
32     {rt_get_node_response, Index::index(), Node::node:node_type()}.

```

The routing table is a `gb_tree`. Identifiers in the ring are integers. Note that in Erlang integer can be of arbitrary precision. For Chord, the identifiers are in $[0, 2^{128})$, i.e. 128-bit strings.

The `rm_beh` behaviour for Chord (excerpt)

File `rt_chord.erl`:

```

44 %% @doc Creates an empty routing table.
45 empty(_Neighbors) -> gb_trees:empty().

```

File `rt_chord.erl`:

```

273 empty_ext(_Neighbors) -> gb_trees:empty().

```

`empty/1` returns an empty `gb_tree`, same for `empty_ext/1`.

`rt_chord:hash_key/1`, `rt_chord:get_random_node_id/0`, `rt_chord:get_replica_keys/1` and `rt_chord:n/0` are implemented like their counterparts in `rt_simple.erl`.

File `rt_chord.erl`:

```

277 %% @doc Returns the next hop to contact for a lookup.
278 %%     If the routing table has less entries than the rt_size_use_neighbors
279 %%     config parameter, the neighborhood is also searched in order to find a
280 %%     proper next hop.
281 %%     Note, that this code will be called from the dht_node process and
282 %%     it will thus have an external_rt!
283 next_hop(State, Id) ->
284     case intervals:in(Id, dht_node_state:get(State, succ_range)) of
285     true -> dht_node_state:get(State, succ_pid);
286     _ ->
287         % check routing table:
288         RT = dht_node_state:get(State, rt),
289         RTSize = get_size(RT),
290         NodeRT = case util:gb_trees_largest_smaller_than(Id, RT) of
291             {value, _Key, N} ->
292                 N;
293             nil when RTSize == 0 ->
294                 dht_node_state:get(State, succ);
295             nil -> % forward to largest finger
296                 {_Key, N} = gb_trees:largest(RT),
297                 N
298         end,
299         FinalNode =
300             case RTSize < config:read(rt_size_use_neighbors) of
301             false -> NodeRT;
302             _ ->
303                 % check neighborhood:
304                 nodelist:largest_smaller_than(
305                     dht_node_state:get(State, neighbors), Id, NodeRT)
306             end,
307         node:pidX(FinalNode)
308     end.

```

If the (external) routing table contains at least one item, the next hop is retrieved from the `gb_tree`. It will be the node with the largest id that is smaller than the id we are looking for. If the routing

table is empty, the successor is chosen. However, if we haven't found the key in our routing table, the next hop will be our largest finger, i.e. entry.

File `rt_chord.erl`:

```

77 %% @doc Starts the stabilization routine.
78 init_stabilize(Neighbors, RT) ->
79     % calculate the longest finger
80     Id = nodelist:nodeid(Neighbors),
81     Key = calculateKey(Id, first_index()),
82     % trigger a lookup for Key
83     lookup:unreliable_lookup(Key, {send_to_group_member, routing_table,
84                                   {rt_get_node, comm:this(), first_index()}}),
85     RT.

```

The routing table stabilization is triggered for the first index and then runs asynchronously, as we do not want to block the `rt_loop` to perform other request while recalculating the routing table.

We have to find the node responsible for the calculated finger and therefore perform a lookup for the node with a `rt_get_node` message, including a reference to ourselves as the reply-to address and the index to be set.

The lookup performs an overlay routing by passing the message until the responsible node is found. There, the message is delivered to the `routing_table` process. The remote node sends the requested information back directly. It includes a reference to itself in a `rt_get_node_response` message. Both messages are handled by `rt_chord:handle_custom_message/2`:

File `rt_chord.erl`:

```

218 %% @doc Chord reacts on 'rt_get_node_response' messages in response to its
219 %%      'rt_get_node' messages.
220 -spec handle_custom_message
221       (custom_message(), rt_loop:state_active()) -> rt_loop:state_active();
222       (any(), rt_loop:state_active()) -> unknown_event.
223 handle_custom_message({rt_get_node, Source_PID, Index}, State) ->
224     MyNode = nodelist:node(rt_loop:get_neighb(State)),
225     comm:send(Source_PID, {rt_get_node_response, Index, MyNode}),
226     State;
227 handle_custom_message({rt_get_node_response, Index, Node}, State) ->
228     OldRT = rt_loop:get_rt(State),
229     Id = rt_loop:get_id(State),
230     Succ = rt_loop:get_succ(State),
231     NewRT = stabilize(Id, Succ, OldRT, Index, Node),
232     check(OldRT, NewRT, rt_loop:get_neighb(State), true),
233     rt_loop:set_rt(State, NewRT);
234 handle_custom_message(_Message, _State) ->
235     unknown_event.

```

File `rt_chord.erl`:

```

150 %% @doc Updates one entry in the routing table and triggers the next update.
151 -spec stabilize(MyId::key() | key_t(), Succ::node:node_type(), OldRT::rt(),
152               Index::index(), Node::node:node_type()) -> NewRT::rt().
153 stabilize(Id, Succ, RT, Index, Node) ->
154     case (node:id(Succ) /= node:id(Node)) % reached succ?
155     andalso (not intervals:in( % there should be nothing shorter
156                               node:id(Node), % than succ
157                               node:mk_interval_between_ids(Id, node:id(Succ)))) of
158     true ->
159         NewRT = gb_trees:enter(Index, Node, RT),
160         Key = calculateKey(Id, next_index(Index)),
161         Msg = {rt_get_node, comm:this(), next_index(Index)},
162         lookup:unreliable_lookup(Key,
163                                   {send_to_group_member, routing_table, Msg}),
164         NewRT;
165     _ -> RT

```

166 end.

stabilize/5 assigns the received routing table entry and triggers the routing table stabilization for the the next shorter entry using the same mechanisms as described above.

If the shortest finger is the successor, then filling the routing table is stopped, as no further new entries would occur. It is not necessary, that Index reaches 1 to make that happen. If less than 2^{128} nodes participate in the system, it may happen earlier.

File `rt_chord.erl`:

```
170 %% @doc Updates the routing table due to a changed node ID, pred and/or succ.
171 -spec update(OldRT::rt(), OldNeighbors::nodelist:neighborhood(),
172             NewNeighbors::nodelist:neighborhood()) -> {trigger_rebuild, rt()}.
173 update(_OldRT, _OldNeighbors, NewNeighbors) ->
174     % to be on the safe side ...
175     {trigger_rebuild, empty(NewNeighbors)}.
```

Tells the `rt_loop` process to rebuild the routing table starting with an empty (internal) routing table state.

File `rt_chord.erl`:

```
89 %% @doc Removes dead nodes from the routing table.
90 filter_dead_node(RT, DeadPid) ->
91     DeadIndices = [Index || {Index, Node} <- gb_trees:to_list(RT),
92                          node:same_process(Node, DeadPid)],
93     lists:foldl(fun(Index, Tree) -> gb_trees:delete(Index, Tree) end,
94                 RT, DeadIndices).
```

`filter_dead_node` removes dead entries from the `gb_tree`.

File `rt_chord.erl`:

```
312 export_rt_to_dht_node(RT, Neighbors) ->
313     Id = nodelist:nodeid(Neighbors),
314     Pred = nodelist:pred(Neighbors),
315     Succ = nodelist:succ(Neighbors),
316     Tree = gb_trees:enter(node:id(Succ), Succ,
317                          gb_trees:enter(node:id(Pred), Pred, gb_trees:empty())),
318     util:gb_trees_foldl(fun (_K, V, Acc) ->
319                         % only store the ring id and the according node structure
320                         case node:id(V) == Id of
321                             true -> Acc;
322                             false -> gb_trees:enter(node:id(V), V, Acc)
323                         end
324     end, Tree, RT).
```

`export_rt_to_dht_node` converts the internal `gb_tree` structure based on indices into the external representation optimised for look-ups, i.e. a `gb_tree` with node ids and the nodes themselves.

File `rt_chord.erl`:

```
239 %% @doc Notifies the dht_node and failure detector if the routing table changed.
240 %%     Provided for convenience (see check/5).
241 check(OldRT, NewRT, Neighbors, ReportToFD) ->
242     check(OldRT, NewRT, Neighbors, Neighbors, ReportToFD).
243
244 %% @doc Notifies the dht_node if the (external) routing table changed.
245 %%     Also updates the failure detector if ReportToFD is set.
246 %%     Note: the external routing table also changes if the Pred or Succ
247 %%     change.
248 check(OldRT, NewRT, OldNeighbors, NewNeighbors, ReportToFD) ->
249     case OldRT == NewRT andalso
```



```

250         nodelist:pred(OldNeighbors) == nodelist:pred(NewNeighbors) andalso
251         nodelist:succ(OldNeighbors) == nodelist:succ(NewNeighbors) of
252     true -> ok;
253 - ->
254     Pid = pid_groups:get_my(dht_node),
255     RT_ext = export_rt_to_dht_node(NewRT, NewNeighbors),
256     comm:send_local(Pid, {rt_update, RT_ext}),
257     % update failure detector:
258     case ReportToFD of
259     true ->
260         NewPids = to_pid_list(NewRT),
261         OldPids = to_pid_list(OldRT),
262         fd:update_subscriptions(OldPids, NewPids);
263     - -> ok
264     end
265 end.

```

Checks whether the routing table changed and in this case sends the `dht_node` an updated (external) routing table state. Optionally the failure detector is updated. This may not be necessary, e.g. if check is called after a crashed node has been reported by the failure detector (the failure detector already unsubscribes the node in this case).

8.4. Local Datastore

8.5. Cyclon

8.6. Vivaldi Coordinates

8.7. Estimated Global Information (Gossiping)

8.8. Load Balancing

8.9. Broadcast Trees

9. Transactions in Scalaris

9.1. The Paxos Module

9.2. Transactions using Paxos Commit

9.3. Applying the Tx-Modules to replicated DHTs

Introduces transaction processing on top of a Overlay

10. How a node joins the system

Description is based on SVN revision r1327.

After booting a new Scalaris-System as described in Section 2.5.1 on page 10, ten additional local nodes can be started by typing `admin:add_nodes(10)` in the Erlang-Shell that the boot process opened ¹.

```
scalaris> ./bin/boot.sh
[...]  
(boot@csr-pc9)1> admin:add_nodes(10)
```

In the following we will trace what this function does in order to add additional nodes to the system. The function `admin:add_nodes(int)` is defined as follows.

File `admin.erl`:

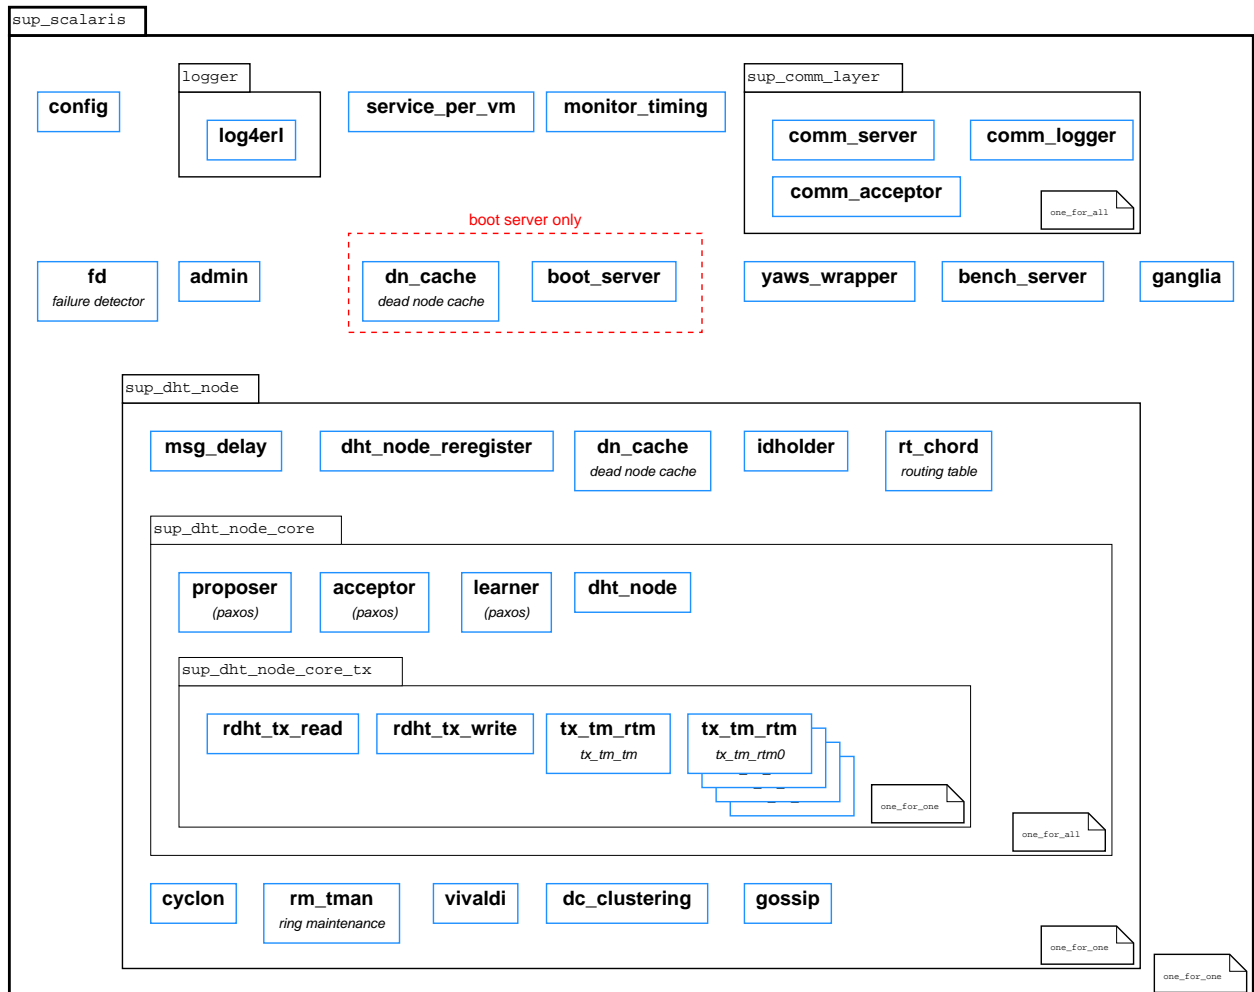
```
37 % @doc add new Scalaris nodes on the local node  
38 -spec add_node_at_id(?RT:key()) ->  
39     ok | {error, already_present} | {already_started, pid() | undefined} | term().  
40 add_node_at_id(Id) ->  
41     add_node([{{idholder, id}, Id}, {skip_psv_lb}]).  
42  
43 -spec add_node([tuple()]) ->  
44     ok | {error, already_present} | {already_started, pid() | undefined} | term().  
45 add_node(Options) ->  
46     DhtNodeId = randoms:getRandomId(),  
47     Desc = util:sup_supervisor_desc(  
48         DhtNodeId, config:read(dht_node_sup), start_link,  
49         [{my_sup_dht_node_id, DhtNodeId} | Options]),  
50     case supervisor:start_child(main_sup, Desc) of  
51         {ok, _Child} -> ok;  
52         {ok, _Child, _Info} -> ok;  
53         {error, _Error} = X -> X  
54     end.  
55  
56 -spec add_nodes(non_neg_integer()) ->  
57     nothing_to_do | [ok | {error, already_present} |  
58         {already_started, pid() | undefined} | term()],...].  
59 add_nodes(0) -> nothing_to_do;  
60 add_nodes(Count) ->  
61     [add_node([]) || _X <- lists:seq(1, Count)].
```

It calls `admin:add_node([])` Count times. This function starts a new child with the given options for the main supervisor `main_sup`. In particular, it sets a random ID that is passed to the new node as its suggested ID to join at. To actually perform the start, the function `sup_dht_node:start_link/1` is called by the Erlang supervisor mechanism. For more details on the OTP supervisor mechanism see Chapter 18 of the Erlang book [1] or the online documentation at <http://www.erlang.org/doc/man/supervisor.html>.

¹Increase the log level to `info` to get more detailed startup logs. See Section 2.7 on page 12

10.1. Supervisor-tree of a Scalaris node

When a new Erlang VM with a Scalaris node is started, a `sup_scalaris` supervisor is started that creates further workers and supervisors according to the following scheme (processes starting order: left to right, top to bottom):



When new nodes are started using `admin:add_node/1`, only new `sup_dht_node` supervisors are started.

10.2. Starting the `sup_dht_node` supervisor and general processes of a node

Starting supervisors is a two step process: a call to `supervisor:start_link/2,3`, e.g. from a custom supervisor's own `start_link` method, will start the supervisor process. It will then call `Module:init/1` to find out about the restart strategy, maximum restart frequency and child processes. Note that `supervisor:start_link/2,3` will not return until `Module:init/1` has returned and all child processes have been started.

Let's have a look at `sup_dht_node:init/1`, the 'DHT node supervisor'.

File `sup_dht_node.erl`:

```

44 -spec init([tuple()]) -> {ok, {{one_for_one, MaxRetries::pos_integer(),
45                               PeriodInSeconds::pos_integer()},
46                               [ProcessDescr::any()]}},
47 init(Options) ->
48   DHTNodeGroup = pid_groups:new("dht_node_"),
49   pid_groups:join_as(DHTNodeGroup, ?MODULE),
50   boot_server:connect(),
51
52   Cyclon = util:sup_worker_desc(cyclon, cyclon, start_link, [DHTNodeGroup]),
53   DC_Clustering =
54     util:sup_worker_desc(dc_clustering, dc_clustering, start_link,
55                           [DHTNodeGroup]),
56   DeadNodeCache =
57     util:sup_worker_desc(deadnodecache, dn_cache, start_link,
58                           [DHTNodeGroup]),
59   Delayer =
60     util:sup_worker_desc(msg_delay, msg_delay, start_link,
61                           [DHTNodeGroup]),
62   Gossip =
63     util:sup_worker_desc(gossip, gossip, start_link, [DHTNodeGroup]),
64   IdHolder =
65     util:sup_worker_desc(idholder, idholder, start_link,
66                           [DHTNodeGroup, Options]),
67   Reregister =
68     util:sup_worker_desc(dht_node_reregister, dht_node_reregister,
69                           start_link, [DHTNodeGroup]),
70   RingMaintenance =
71     util:sup_worker_desc(ring_maintenance, rm_loop, start_link, [DHTNodeGroup]),
72   RoutingTable =
73     util:sup_worker_desc(routing_table, rt_loop, start_link,
74                           [DHTNodeGroup]),
75   SupDHTNodeCore_AND =
76     util:sup_supervisor_desc(sup_dht_node_core, sup_dht_node_core,
77                               start_link, [DHTNodeGroup, Options]),
78   Vivaldi =
79     util:sup_worker_desc(vivaldi, vivaldi, start_link, [DHTNodeGroup]),
80
81   %% order in the following list is the start order
82   {ok, {{one_for_one, 10, 1},
83         [
84           Delayer,
85           Reregister,
86           DeadNodeCache,
87           IdHolder,
88           RingMaintenance,
89           RoutingTable,
90           Cyclon,
91           Vivaldi,
92           DC_Clustering,
93           Gossip,
94           SupDHTNodeCore_AND
95         ]}}}.

```

The return value of the `init/1` function specifies the child processes of the supervisor and how to start them. Here, we define a list of processes to be observed by a `one_for_one` supervisor. The processes are: `Delayer`, `Reregister`, `DeadNodeCache`, `IdHolder`, `RoutingTable`, `SupDHTNodeCore_AND`, `Cyclon`, `RingMaintenance`, `Vivaldi`, `DC_Clustering` and a `Gossip` process in this order.

The term `{one_for_one, 10, 1}` specifies that the supervisor should try 10 times to restart each process before giving up. `one_for_one` supervision means, that if a single process stops, only that process is restarted. The other processes run independently.

When the `sup_dht_node:init/1` is finished the supervisor module starts all the defined processes by calling the functions that were defined in the returned list.

For a join of a new node, we are only interested in the starting of the `SupDHTNodeCore_AND` process here. At that point in time, all other defined processes are already started and running.

10.3. Starting the `sup_dht_node_core` supervisor with a peer and some paxos processes

Like any other supervisor the `sup_dht_node_core` supervisor calls its `sup_dht_node_core:init/1` function:

File `sup_dht_node_core.erl`:

```
40 -spec init({pid_groups:groupname(), Options::[tuple()]}) ->
41     {ok, {{one_for_all, MaxRetries::pos_integer(),
42             PeriodInSeconds::pos_integer()},
43           [ProcessDescr::any()]}}.
44 init({DHTNodeGroup, Options}) ->
45     pid_groups:join_as(DHTNodeGroup, ?MODULE),
46     Proposer =
47         util:sup_worker_desc(proposer, proposer, start_link, [DHTNodeGroup]),
48     Acceptor =
49         util:sup_worker_desc(acceptor, acceptor, start_link, [DHTNodeGroup]),
50     Learner =
51         util:sup_worker_desc(learner, learner, start_link, [DHTNodeGroup]),
52     DHTNode =
53         util:sup_worker_desc(dht_node, dht_node, start_link,
54                             [DHTNodeGroup, Options]),
55     TX =
56         util:sup_supervisor_desc(sup_dht_node_core_tx, sup_dht_node_core_tx, start_link,
57                                 [DHTNodeGroup]),
58     {ok, {{one_for_all, 10, 1},
59         [
60             Proposer, Acceptor, Learner,
61             DHTNode,
62             TX
63         ]}}.
```

It defines five processes, that have to be observed using a `one_for_all`-supervisor, which means, that if one fails, all have to be restarted. The `dht_node` module implements the main component of a full Scalaris node which glues together all the other processes. Its `dht_node:start_link/2` function will get the following parameters: (a) the processes' group that is used with the `pid_groups` module and (b) a list of options for the `dht_node`. The process group name was calculated a bit earlier in the code. *Exercise: Try to find where.*

File `dht_node.erl`:

```
395 %% @doc spawns a scalaris node, called by the scalaris supervisor process
396 -spec start_link(pid_groups:groupname(), [tuple()]) -> {ok, pid()}.
397 start_link(DHTNodeGroup, Options) ->
398     gen_component:start_link(?MODULE, Options,
399                             [{pid_groups_join_as, DHTNodeGroup, dht_node}, wait_for_init]).
```

Like many other modules, the `dht_node` module implements the `gen_component` behaviour. This behaviour was developed by us to enable us to write code which is similar in syntax and semantics to the examples in [3]. Similar to the supervisor behaviour, a module implementing this behaviour has to provide an `init/1` function, but here it is used to initialize the state of the component. This function is described in the next section.

Note: `?MODULE` is a predefined Erlang macro, which expands to the module name, the code belongs to (here: `dht_node`).

10.4. Initializing a `dht_node-process`

File dht_node.erl:

```
375 %% @doc joins this node in the ring and calls the main loop
376 -spec init(Options::[tuple()]) -> {join, {as_first | phase1, Options::[tuple()]},
377                                     msg_queue:msg_queue()}.
378 init(Options) ->
379     {my_sup_dht_node_id, MySupDhtNode} = lists:keyfind(my_sup_dht_node_id, 1, Options),
380     erlang:put(my_sup_dht_node_id, MySupDhtNode),
381     % first node in this vm and also vm is marked as first
382     % or unit-test
383     case is_first(Options) of
384         true ->
385             trigger_known_nodes(),
386             idholder:get_id(),
387             {join, {as_first, Options}, msg_queue:new()};
388         _ ->
389             idholder:get_id(),
390             {join, {phase1, Options}, msg_queue:new()}
391     end.
```

The `gen_component` behaviour registers the `dht_node` in the process dictionary. Formerly, the process had to do this itself, but we moved this code into the behaviour. If the `dht_node` is the first node, it will start immediately by triggering all known nodes (to initialize the comm layer) and entering the join process accordingly. The node also retrieves its `Id` from the `idholder`: `idholder:get_id()`. In case of a re-started node, this will be the last known ID. Otherwise, the `idholder` will return the ID (first) set by the `{idholder, id}, KEY` option passed to `idholder:init/1` or a random key. If the `dht_node`-process fails and is re-started by its supervisor, this call to the `idholder` ensures that the node still keeps its `Id`, assuming that the `idholder` process is not failing and no passive load balancing takes place. This is important in order not to destroy an existing load balance if a node somehow fails.

10.5. Actually joining the ring

After retrieving its identifier, the node starts the join protocol which processes the appropriate messages calling `dht_node_join:process_join_msg(Message, State)`.

10.5.1. A single node joining an empty ring

File dht_node_join.erl:

```
77 process_join_state({idholder_get_id_response, Id, IdVersion},
78                    {join, {as_first, _Options}, QueuedMessages}) ->
79     log:log(info, "[ Node ~w ] joining as first: ~p", [self(), Id]),
80     Me = node:new(comm:this(), Id, IdVersion),
81     % join complete, State is the first "State"
82     finish_join(Me, Me, Me, ?DB:new(), QueuedMessages);
```

If the ring is empty, the joining node will be the only node in the ring and will thus be responsible for the whole key space. `dht_node_join:finish_join/5` just creates a new state for a `Scalaris` node consisting of the given parameters (the node as itself, its predecessor and successor, an empty database and the queued messages that arrived during the join). It then activates all dependent processes and creates a routing table from this information.

The `dht_node_state:state()` type is defined in

File dht_node_state.erl:

```

50 -record(state, {rt           = ?required(state, rt)           :: ?RT:external_rt(),
51                  neighbors   = ?required(state, neighbors)  :: tid(),
52                  join_time    = ?required(state, join_time)  :: util:time(),
53                  trans_log    = ?required(state, trans_log)  :: #translog{},
54                  db           = ?required(state, db)         :: ?DB:db(),
55                  tx_tp_db     = ?required(state, tx_tp_db)   :: any(),
56                  proposer     = ?required(state, proposer)  :: pid(),
57                  % slide with pred (must not overlap with 'slide with succ!':
58                  slide_pred = null :: slide_op:slide_op() | null,
59                  % slide with succ (must not overlap with 'slide with pred!':
60                  slide_succ = null :: slide_op:slide_op() | null,
61                  msg_fwd     = [] :: [{intervals:interval(), comm:mypid()}],
62                  % additional range to respond to during a move:
63                  db_range     = intervals:empty() :: intervals:interval()
64                  }).
65 -opaque state() :: #state{}.

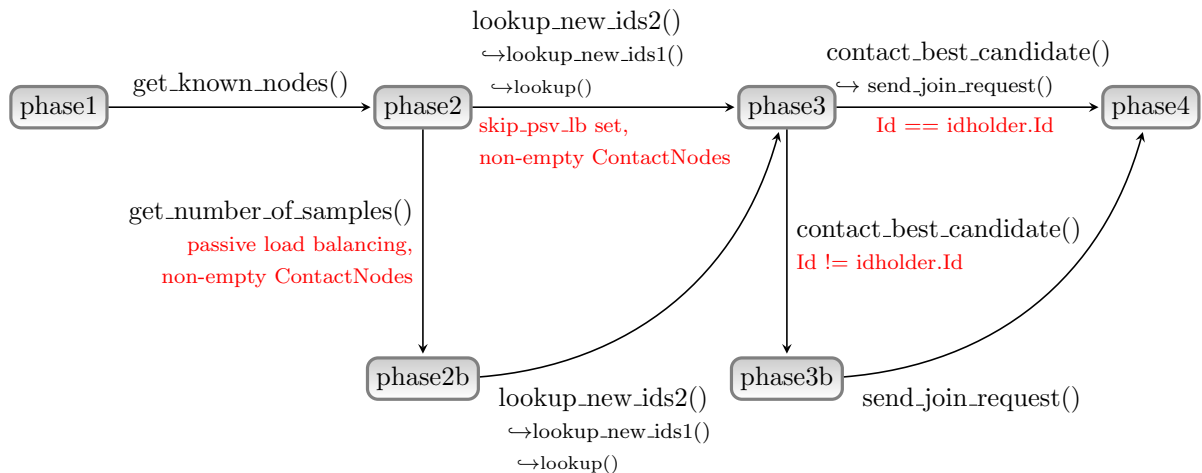
```

10.5.2. A single node joining an existing (non-empty) ring

If a node joins an existing ring, its join protocol will step through the following 6 phases:

- **phase1** getting the Id from the idholder
- **phase2** finding nodes to contact with the help of the configured `known_hosts`
- **phase2b** getting the number of Ids to sample (may be skipped)
- **phase3** lookup nodes responsible for all sampled Ids
- **phase3b** waiting for the idholder to acknowledge an Id change (if applicable)
- **phase4** joining a selected node and setting up item movements

The following figure shows a (non-exhaustive) overview of the transitions between the phases in the normal case. We will go through these step by step and discuss what happens if errors occur.



Phase 1

Phase 1 is entered during the initialization of the `dht_node` process after requesting the ID from the idholder. It receives its `idholder_get_id_response` message and contacts all nodes set in the `known_hosts` configuration parameter. Their responses are handled in phase 2.

File dht_node_join.erl:

```
89 % 1. get my key
90 process_join_state({idholder_get_id_response, Id, IdVersion},
91                   {join, {phase1, Options}, QueuedMessages}) ->
92     %io:format("p1: got key~n"),
93     log:log(info, "[ Node ~w ] joining", [self()]),
94     get_known_nodes(),
95     msg_delay:send_local(get_join_timeout() div 1000, self(), {join, timeout}),
96     {join, {phase2, Options, Id, IdVersion, [], [Id], []}, QueuedMessages};
```

Phase 2 and 2b

Phase 2 collects all dht_node processes inside the contacted VMs. It therefor mainly processes get_dht_nodes_response messages and integrates all received nodes into the list of ContactNodes. The next step depends on whether the {skip_psv_lb} option for skipping any passive load balancing algorithm has been given to the dht_node or not. If it is present, the node will use the ID from the idholder, issue a lookup for the responsible node and move to phase 3. Otherwise, the passive load balancing's lb_psv_*:get_number_of_samples/1 method will be called asking for the number of IDs to sample. Its answer will be processed in phase 2b.

get_dht_nodes_response messages arriving in phase 2b will be processed the same way since no response for the request for the number of nodes to sample has been processed yet. For better efficiency, get_dht_nodes_response messages arriving in any later phase will be processed anyway and the send dht_node processes will be integrated into the ContactNodes. These phases' operations will however not be interrupted and nothing else is changed.

File dht_node_join.erl:

```
111 % in phase 2, 2b add the nodes and do lookups with them / get number of samples
112 process_join_state({get_dht_nodes_response, Nodes},
113                   {join, JoinState, QueuedMessages})
114     when element(1, JoinState) == phase2 orelse
115         element(1, JoinState) == phase2b ->
116         %io:format("p2: got dht_nodes_response ~p~n", [lists:delete(comm:this(), Nodes)]),
117         ContactNodes = [Node || Node <- Nodes, Node /= comm:this()],
118         NewJoinState =
119             case skip_psv_lb(JoinState) of
120             true -> % skip phase2b (use only the given ID)
121                 JoinState1 = add_contact_nodes_back(ContactNodes, JoinState),
122                 % (re-)issue lookups for all existing IDs
123                 JoinState2 = lookup(JoinState1),
124                 % the Id may have been removed -> select a new one in this case:
125                 lookup_new_ids2(1, JoinState2);
126             _ -> get_number_of_samples(JoinState, ContactNodes, true)
127             end,
128         {join, NewJoinState, QueuedMessages};
129
130 % in all other phases, just add the provided nodes:
131 % note: phase1 should never receive this message!
132 process_join_state({get_dht_nodes_response, Nodes},
133                   {join, JoinState, QueuedMessages})
134     when element(1, JoinState) == phase3 orelse
135         element(1, JoinState) == phase3b orelse
136         element(1, JoinState) == phase4 ->
137         FurtherNodes = [Node || Node <- Nodes, Node /= comm:this()],
138         {join, add_contact_nodes_back(FurtherNodes, JoinState), QueuedMessages};
```

Phase 2b will handle get_number_of_samples messages from the passive load balance algorithm. Once received, new (unique) IDs will be sampled randomly so that the total number of join candidates (selected IDs together with fully processed candidates from further phases) is at least as high

as the given number of samples. Afterwards, lookups will be created for all previous IDs as well as the new ones and the node will move to phase 3.

File dht_node_join.erl:

```
154 % note: although this message was send in phase2, also accept message in
155 % phase2, e.g. messages arriving from previous calls
156 process_join_state({join, get_number_of_samples, Samples, Source},
157                   {join, JoinState, QueuedMessages})
158   when element(1, JoinState) == phase2 orelse
159     element(1, JoinState) == phase2b ->
160     %io:format("p2b: got get_number_of_samples ~p~n", [Samples]),
161     % prefer node that send get_number_of_samples as first contact node
162     JoinState1 = add_contact_nodes_front([Source], JoinState),
163     % (re-)issue lookups for all existing IDs
164     JoinState2 = lookup(JoinState1),
165     % then create additional samples, if required
166     {join, lookup_new_ids2(Samples, JoinState2), QueuedMessages};
167
168 % ignore message arriving in other phases:
169 process_join_state({join, get_number_of_samples, _Samples, _Source}, State) ->
170   State;
```

Lookups will make Scalaris find the node currently responsible for a given ID and send a request to simulate a join to this node, i.e. a get_candidate message. Note that during such an operation, the joining node will become the existing node's predecessor. The simulation will be delegated to the passive load balance algorithm the joining node requested, as set by the join_lb_psv configuration parameter.

File dht_node_join.erl:

```
371 process_join_msg({join, get_candidate, Source_PID, Key, LbPsv}, State) ->
372   % if anything goes wrong creating the candidate, do not crash, just report
373   % a no_op operation to the other node
374   LbPsv:create_join(State, Key, Source_PID);
```

Phase 3

The result of the simulation will be send in a get_candidate_response message and will be processed in phase 3 of the joining node. It will be integrated into the list of processed candidates. If there are no more IDs left to process, the best among them will be contacted. Otherwise further get_candidate_response messages will be awaited. Such messages will also be processed in the other phases where the candidate will be simply added to the list.

File dht_node_join.erl:

```
186 process_join_state({join, get_candidate_response, OrigJoinId, Candidate},
187                   {join, JoinState, QueuedMessages})
188   when element(1, JoinState) == phase3 ->
189     %io:format("p3: lookup success~n"),
190     JoinState1 = remove_join_id(OrigJoinId, JoinState),
191     JoinState2 = integrate_candidate(Candidate, JoinState1, front),
192     NewJoinState =
193       case get_join_ids(JoinState2) of
194         [] -> % no more join ids to look up -> join with the best:
195           contact_best_candidate(JoinState2);
196         [_|_] -> % still some unprocessed join ids -> wait
197           JoinState2
198       end,
199     {join, NewJoinState, QueuedMessages};
200
201 % In phase 2 or 2b, also add the candidate but do not continue.
```

```

202 process_join_state({join, get_candidate_response, OrigJoinId, Candidate},
203                    {join, JoinState, QueuedMessages})
204   when element(1, JoinState) == phase2 orelse
205       element(1, JoinState) == phase2b ->
206       JoinState1 = remove_join_id(OrigJoinId, JoinState),
207       JoinState2 = integrate_candidate(Candidate, JoinState1, front),
208       {join, JoinState2, QueuedMessages};
209
210 % In phase 3b or 4, add the candidate to the end of the candidates as they are sorted
211 % and the join with the first has already started (use this candidate as backup
212 % if the join fails). Do not start a new join.
213 process_join_state({join, get_candidate_response, OrigJoinId, Candidate},
214                    {join, JoinState, QueuedMessages})
215   when element(1, JoinState) == phase3b orelse
216       element(1, JoinState) == phase4 ->
217       JoinState1 = remove_join_id(OrigJoinId, JoinState),
218       JoinState2 = integrate_candidate(Candidate, JoinState1, back),
219       {join, JoinState2, QueuedMessages};

```

If `dht_node_join:contact_best_candidate/1` is called and candidates are available (there should be at this stage!), it will sort the candidates by using the passive load balance algorithm and get the ID from the best candidate. Note that during simulation at the existing node, this ID may be different from the initial ID send during the lookup. If this ID is the same as the one set at the `idholder`, we can continue directly by sending a `join_request` message and going to phase 4.

File `dht_node_join.erl`:

```

606 %% @doc Contacts the best candidate among all stored candidates and sends a
607 %%      join_request (Timeouts = 0).
608 -spec contact_best_candidate(JoinState::phase_2_4())
609       -> phase2() | phase2b() | phase3b() | phase4().
610 contact_best_candidate(JoinState) ->
611     contact_best_candidate(JoinState, 0).
612 %% @doc Contacts the best candidate among all stored candidates and sends a
613 %%      join_request. Timeouts is the number of join_request_timeout messages
614 %%      previously received.
615 -spec contact_best_candidate(JoinState::phase_2_4(), Timeouts::non_neg_integer())
616       -> phase2() | phase2b() | phase3b() | phase4().
617 contact_best_candidate(JoinState, Timeouts) ->
618     JoinState1 = sort_candidates(JoinState),
619     case get_candidates(JoinState1) of
620     [] -> % no candidates -> start over:
621         start_over(JoinState1);
622     [BestCand | _] ->
623         NewId = node_details:get(lb_op:get(BestCand, n1_new), new_key),
624         Id = get_id(JoinState1),
625         IdVersion = get_id_version(JoinState1),
626         case NewId of
627         Id -> send_join_request(JoinState1, Timeouts);
628         _ -> idholder:set_id(NewId, IdVersion + 1),
629             idholder:get_id(),
630             set_phase(phase3b, JoinState1)
631         end
632     end.

```

File `dht_node_join.erl`:

```

636 %% @doc Sends a join request to the first candidate. Timeouts is the number of
637 %%      join_request_timeout messages previously received.
638 %%      PreCond: the id has been set to the ID to join at and has been updated
639 %%      in JoinState.
640 -spec send_join_request(JoinState::phase_2_4(), Timeouts::non_neg_integer())
641       -> phase2() | phase2b() | phase4().
642 send_join_request(JoinState, Timeouts) ->
643     case get_candidates(JoinState) of
644     [] -> % no candidates -> start over (should not happen):
645         start_over(JoinState);

```

```

646     [BestCand | _] ->
647         Id = get_id(JoinState),
648         IdVersion = get_id_version(JoinState),
649         NewSucc = node_details:get(lb_op:get(BestCand, n1succ_new), node),
650         Me = node:new(comm:this(), Id, IdVersion),
651         comm:send(node:pidX(NewSucc), {join, join_request, Me}),
652         msg_delay:send_local(get_join_request_timeout() div 1000,
653                               self(), {join, join_request_timeout, Timeouts}),
654         set_phase(phase4, JoinState)
655     end.

```

Otherwise we need to set the ID at the idholder first and wait for its acknowledgement in phase 3b. A join_request is send once the idholder_get_id_response message arrives.

File dht_node_join.erl:

```

223 process_join_state({idholder_get_id_response, Id, IdVersion},
224                    {join, JoinState, QueuedMessages})
225     when element(1, JoinState) == phase3b->
226         %io:format("p3b: got key~n"),
227         JoinState1 = set_id(Id, IdVersion, JoinState),
228         {join, send_join_request(JoinState1, 0), QueuedMessages};

```

The join_request message will be received by the existing node which will set up a slide operation with the new node or deny the request if it is not responsible for the key (anymore) and reply with a {join, join_response, not_responsible, Node} message.

File dht_node_join.erl:

```

377 process_join_msg({join, join_request, NewPred}, State) when (not is_atom(NewPred)) ->
378     TargetId = node:id(NewPred),
379     % only reply to join request with keys in our range:
380     KeyInRange = dht_node_state:is_responsible(node:id(NewPred), State),
381     case KeyInRange andalso
382         dht_node_move:can_slide_pred(State, TargetId, {join, 'rcv'}) of
383     true ->
384         % TODO: implement step-wise join
385         MoveFullId = util:get_global_uid(),
386         SlideOp = slide_op:new_sending_slide_join(
387             MoveFullId, NewPred, join, State),
388         SlideOp1 = slide_op:set_phase(SlideOp, wait_for_pred_update_join),
389         RMSubscrTag = {move, slide_op:get_id(SlideOp1)},
390         rm_loop:subscribe(self(), RMSubscrTag,
391                           fun(_OldNeighbors, NewNeighbors) ->
392                               NewPred == nodelist:pred(NewNeighbors)
393                           end,
394                           fun dht_node_move:rm_notify_new_pred/4),
395         State1 = dht_node_state:add_db_range(
396             State, slide_op:get_interval(SlideOp1)),
397         send_join_response(State1, SlideOp1, NewPred);
398     _ when not KeyInRange ->
399         comm:send(node:pidX(NewPred), {join, join_response, not_responsible,
400                                         dht_node_state:get(State, node)}),
401         State;
402     _ -> State
403     end;

```

If it is responsible for the ID and is not participating in a slide with its current predecessor, it will set up a slide with the joining node:

File dht_node_join.erl:

```

675 -spec send_join_response(State::dht_node_state:state(),
676                          NewSlideOp::slide_op:slide_op(),
677                          NewPred::node:node_type())
678     -> dht_node_state:state().

```

```

679 send_join_response(State, SlideOp, NewPred) ->
680     MoveFullId = slide_op:get_id(SlideOp),
681     NewSlideOp = slide_op:set_timer(SlideOp, get_join_response_timeout(),
682                                     {join, join_response_timeout, NewPred, MoveFullId}),
683     MyOldPred = dht_node_state:get(State, pred),
684     MyNode = dht_node_state:get(State, node),
685     comm:send(node:pidX(NewPred), {join, join_response, MyNode, MyOldPred, MoveFullId}),
686     % no need to tell the ring maintenance -> the other node will trigger an update
687     % also this is better in case the other node dies during the join
688     %% rm_loop:notify_new_pred(comm:this(), NewPred),
689     dht_node_state:set_slide(State, pred, NewSlideOp).

```

Phase 4

The joining node will receive the `join_response` message in phase 4 of the join protocol. If everything is ok, it will notify its ring maintenance process that it enters the ring, start all required processes and join the slide operation set up by the existing node in order to receive some of its data.

If the join candidate's node is not responsible for the candidate's ID anymore or the candidate's ID already exists, the next candidate is contacted until no further candidates are available and the join protocol starts over using `dht_node_join:start_over/1`.

Note that the `join_response` message will actually be processed in any phase. Therefore, if messages arrive late, the join can be processed immediately and the rest of the join protocol does not need to be executed again.

File `dht_node_join.erl`:

```

253 process_join_state({join, join_response, not_responsible, Node},
254                   {join, JoinState, QueuedMessages} = State)
255     when element(1, JoinState) == phase4 ->
256         % the node we contacted is not responsible for our key (anymore)
257         % -> try the next candidate, if the message is related to the current candidate
258         case get_candidates(JoinState) of
259             [] -> % no candidates -> should not happen in phase4!
260                 log:log(error, "[ Node ~w ] empty candidate list in join phase 4, "
261                               "starting over", [self()]),
262                 {join, start_over(JoinState), QueuedMessages};
263             [Candidate | _Rest] ->
264                 CandidateNode = node_details:get(lb_op:get(Candidate, n1succ_new), node),
265                 case node:same_process(CandidateNode, Node) of
266                     false -> State; % ignore old/unrelated message
267                     _ ->
268                         log:log(info,
269                               "[ Node ~w ] node contacted for join is not responsible for "
270                               "the selected ID (anymore) ~w, trying next candidate",
271                               [self(), get_id(JoinState)]),
272                         JoinState1 = remove_candidate_front(JoinState),
273                         {join, contact_best_candidate(JoinState1), QueuedMessages}
274                 end
275         end;
276     end;
277 % ignore (delayed) message arriving in other phases:
278 process_join_state({join, join_response, not_responsible, _Node}, State) -> State;
279
280 % note: accept (delayed) join_response messages in any phase
281 process_join_state({join, join_response, Succ, Pred, MoveId},
282                   {join, JoinState, QueuedMessages} = State) ->
283     %io:format("p4: join_response~n"),
284     % only act on related messages, i.e. messages from the current candidate
285     Phase = get_phase(JoinState),
286     case get_candidates(JoinState) of
287         [] when Phase == phase4 -> % no candidates -> should not happen in phase4!
288             log:log(error, "[ Node ~w ] empty candidate list in join phase 4, "

```

```

289         "starting over", [self()]),
290         {join, start_over(JoinState), QueuedMessages};
291     [] -> State; % in all other phases, ignore the delayed join_response
292             % if no candidates exist
293     [Candidate | _Rest] ->
294 %%         io:format("Candidate: ~.0p~n", [Candidate]),
295         CandidateNode = node_details:get(lb_op:get(Candidate, n1succ_new), node),
296         case node:same_process(CandidateNode, Succ) of
297             false -> State; % ignore old/unrelated message
298             _ ->
299                 MyId = get_id(JoinState),
300                 MyIdVersion = get_id_version(JoinState),
301                 case MyId == node:id(Succ) orelse MyId == node:id(Pred) of
302                     true ->
303                         log:log(warn, "[ Node ~w ] chosen ID already exists, "
304                             "trying next candidate", [self()]),
305                         JoinState1 = remove_candidate_front(JoinState),
306                         {join, contact_best_candidate(JoinState1), QueuedMessages};
307                     _ ->
308                         Me = node:new(comm:this(), MyId, MyIdVersion),
309                         log:log(info, "[ Node ~w ] joined between ~w and ~w",
310                             [self(), Pred, Succ]),
311                         rm_loop:notify_new_succ(node:pidX(Pred), Me),
312                         rm_loop:notify_new_pred(node:pidX(Succ), Me),
313
314                         finish_join_and_slide(Me, Pred, Succ, ?DB:new(),
315                             QueuedMessages, MoveId)
316
317                 end
318         end;
end;

```

File dht_node_join.erl:

```

693 -spec finish_join(Me::node:node_type(), Pred::node:node_type(),
694                 Succ::node:node_type(), DB::?DB:db(),
695                 QueuedMessages::msg_queue:msg_queue())
696     -> dht_node_state:state().
697 finish_join(Me, Pred, Succ, DB, QueuedMessages) ->
698     rm_loop:activate(Me, Pred, Succ),
699     % wait for the ring maintenance to initialize and tell us its table ID
700     NeighbTable = rm_loop:get_neighbors_table(),
701     rt_loop:activate(NeighbTable),
702     cyclon:activate(),
703     vivaldi:activate(),
704     dc_clustering:activate(),
705     gossip:activate(),
706     dht_node_reregister:activate(),
707     msg_queue:send(QueuedMessages),
708     NewRT_ext = ?RT:empty_ext(rm_loop:get_neighbors(NeighbTable)),
709     dht_node_state:new(NewRT_ext, NeighbTable, DB).
710
711 -spec finish_join_and_slide(Me::node:node_type(), Pred::node:node_type(),
712                             Succ::node:node_type(), DB::?DB:db(),
713                             QueuedMessages::msg_queue:msg_queue(), MoveId::slide_op:id())
714     -> dht_node_state:state().
715 finish_join_and_slide(Me, Pred, Succ, DB, QueuedMessages, MoveId) ->
716     State = finish_join(Me, Pred, Succ, DB, QueuedMessages),
717     SlideOp = slide_op:new_receiving_slide_join(MoveId, Pred, Succ, node:id(Me), join),
718     SlideOp1 = slide_op:set_phase(SlideOp, wait_for_node_update),
719     State1 = dht_node_state:set_slide(State, succ, SlideOp1),
720     State2 = dht_node_state:add_msg_fwd(
721         State1, slide_op:get_interval(SlideOp1),
722         node:pidX(slide_op:get_node(SlideOp1))),
723     RMSubscrTag = {move, slide_op:get_id(SlideOp1)},
724     NewMsgQueue = msg_queue:add(QueuedMessages,
725         {move, node_update, RMSubscrTag}),
726     msg_queue:send(NewMsgQueue),
727     State2.

```

The macro `?RT` maps to the configured routing algorithm. It is defined in `include/scalaris.hrl`. For further details on the routing see Chapter 8.3 on page 31.

Timeouts and other errors

The `join_timeout` parameter in the config files defines an overall timeout for the whole join operation. If it takes longer than `join_timeout` ms, the join will be re-started:

File `dht_node_join.erl`:

```
659 %% @doc Goes back to phase 2 or 2b depending on whether contact nodes are
660 %%      available or not.
661 -spec start_over(JoinState::phase_2_4()) -> phase2() | phase2b().
662 start_over(JoinState) ->
663     case get_contact_nodes(JoinState) of
664         [] ->
665             get_known_nodes(),
666             set_phase(phase2, JoinState);
667         [_|_] = ContactNodes ->
668             get_number_of_samples(JoinState, ContactNodes, false)
669     end.
```

Misc. (all phases)

Note that join-related messages arriving in other phases than those handling them will be ignored. Any other messages during a `dht_node`'s join will be queued and re-send when the join is complete.

11. Directory Structure of the Source Code

The directory tree of Scalaris is structured as follows:

<code>bin</code>	contains shell scripts needed to work with Scalaris (e.g. start the boot services, start a node, ...)
<code>contrib</code>	necessary third party packages (yaws and log4erl)
<code>doc</code>	generated Erlang documentation
<code>docroot</code>	root directory of the node's webserver
<code>ebin</code>	the compiled Erlang code (beam files)
<code>java-api</code>	a Java API to Scalaris
<code>log</code>	log files
<code>src</code>	contains the Scalaris source code
<code>test</code>	unit tests for Scalaris
<code>user-dev-guide</code>	contains the sources for this document

12. Java API

For the Java API documentation, we refer the reader to the documentation generated by javadoc or doxygen. The following commands create the documentation:

```
%> cd java-api  
%> ant doc  
%> doxygen
```

The documentation can then be found in `java-api/doc/index.html` (javadoc) and `java-api/doc-doxygen/html/index.html` (doxygen).

We provide two kinds of APIs:

- high-level access with `de.zib.scalarisc.Scalaris`
- low-level access with `de.zib.scalarisc.Transaction`

The former provides general functions for reading, writing and deleting single key-value pairs and an API for the built-in PubSub-service. The latter allows the user to write custom transactions which can modify an arbitrary number of key-value pairs within one transaction.

Bibliography

- [1] Joe Armstrong. *Programming Erlang: Software for a Concurrent World*. Pragmatic Programmers, ISBN: 978-1-9343560-0-5, July 2007
- [2] Frank Dabek, Russ Cox, Frans Kaashoek, Robert Morris. *Vivaldi: A Decentralized Network Coordinate System*. ACM SIGCOMM 2004.
- [3] Rachid Guerraoui and Luis Rodrigues. *Introduction to Reliable Distributed Programming*. Springer-Verlag, 2006.
- [4] Ion Stoica, Robert Morris, David Karger, M. Frans Kaashoek and Hari Balakrishnan. *Chord: A Scalable Peer-to-peer Lookup Service for Internet Applications*. ACM SIGCOMM 2001, San Deigo, CA, August 2001, pp. 149-160. http://pdos.csail.mit.edu/papers/chord:sigcomm01/chord_sigcomm.pdf
- [5] Mark Jelasity, Alberto Montresor, Özalp Babaoglu. *T-Man: Gossip-based fast overlay topology construction*. Computer Networks (CN) 53(13):2321-2339, 2009.
- [6] F. Schintke, A. Reinefeld, S. Haridi, T. Schütt. *Enhanced Paxos Commit for Transactions on DHTs*. 10th IEEE/ACM Int. Conf. on Cluster, Cloud and Grid Computing, pp. 448-454, May 2010.
- [7] Spyros Voulgaris, Daniela Gavidia, Maarten van Steen. *CYCLON: Inexpensive Membership Management for Unstructured P2P Overlays*. J. Network Syst. Manage. 13(2): 2005.
- [8] Márk Jelasity, Alberto Montresor, Ozalp Babaoglu. *Gossip-based aggregation in large dynamic networks*. ACM Trans. Comput. Syst. 23(3), 219-252 (2005).

Index

?RT
 next_hop, 32
 update, 34

admin
 add_node, 43, 44

comm, 3, 23, 23
 get_msg_tag, 27
 send_to_group_member, 27

cs_api, 24

dht_node, 33, 34, 37, 41, 46, 48, 49
dht_node_join
 contact_best_candidate, 51, 51
 finish_join, 47, 54
 finish_join_and_slide, 54
 process_join_msg, 47
 send_join_request, 51
 send_join_response, 52
 start_over, 53, 55

dht_node_state
 state, 47

erlang
 exit, 24, 26
 now, 21
 send_after, 21

ets
 i, 21

gen_component, 3, 21, 23, 23–30
 bp_barrier, 28
 bp_cont, 28, 29
 bp_del, 27, 28
 bp_set, 27
 bp_set_cond, 27
 bp_step, 28–30
 change_handler, 24, 26, 26
 get_state, 24
 kill, 25, 26
 runnable, 28
 sleep, 26
 start, 25
 start_link, 25, 26

idholder, 47–49, 51, 52
 get_id, 47
 init, 47

intervals
 in, 32

lb_psv_*
 get_number_of_samples, 49

msg_delay, 21

paxos_SUITE, 27
 step_until_decide, 29

pdb, 30

pid_groups, 3, 23, 23, 25, 27, 46

randoms, 35

rm_beh, 35, 38

routing_table, 39

rt_beh, 31
 check, 33
 check_config, 33
 dump, 33
 empty, 32
 empty_ext, 33
 export_rt_to_dht_node, 33
 filter_dead_node, 33
 get_random_node_id, 33
 get_replica_keys, 33
 get_size, 33
 handle_custom_message, 33
 hash_key, 33
 init_stabilize, 33
 n, 33
 next_hop, 33
 to_list, 33
 to_pid_list, 33
 update, 33

rt_chord, 37
 empty, 38
 empty_ext, 38
 export_rt_to_dht_node, 40

- filter_dead_node, [40](#)
- get_random_node_id, [38](#)
- get_replica_keys, [38](#)
- handle_custom_message, [39](#), [39](#)
- hash_key, [38](#)
- init_stabilize, [39](#)
- n, [38](#)
- next_hop, [38](#)
- stabilize, [39](#)
- update, [40](#)

rt_loop, [33](#), [33](#), [40](#)

rt_simple, [34](#)

- dump, [36](#)
- empty, [35](#)
- empty_ext, [35](#)
- export_rt_to_dht_node, [36](#)
- filter_dead_node, [36](#)
- get_random_node_id, [35](#)
- get_replica_keys, [36](#)
- get_size, [36](#)
- handle_custom_message, [37](#)
- hash_key, [35](#)
- init_stabilize, [35](#)
- n, [36](#)
- next_hop, [35](#)
- to_list, [36](#)
- to_pid_list, [36](#)
- update, [35](#)

sup_dht_node

- init, [44](#), [45](#)
- start_link, [43](#)

sup_dht_node_core, [46](#)

sup_scalaris, [44](#)

supervisor

- start_link, [44](#)

timer

- sleep, [24](#)
- tc, [21](#)

util

- tc, [21](#)

vivaldi, [27](#)

vivaldi_latency, [27](#)

your_gen_component

- init, [24](#), [26](#)
- on, [24–26](#)